

Introduction to Computer Science

Jürgen Schönwälder

Jacobs University Bremen

November 27, 2017



JACOBS
UNIVERSITY



<http://cnds.eecs.jacobs-university.de/courses/ics-2017>

Part: Administrivia

1 Course Objectives and Grading

2 Rules of the Game

3 Resources

Course Objectives and Grading

1 Course Objectives and Grading

2 Rules of the Game

3 Resources

Topics and learning goals

- Become familiar with core concepts of computer science
- Learn key mathematical foundations (sets, relations, functions, proofs, ...)
- Understand boolean logic and how it relates to digital circuits
- Understand basics of computer architecture and system software
- Learn about abstract machines (finite state, push-down, turing machines)
- Understand which problems can (not) be solved by a computer
- Learn different programming paradigms (procedural, functional, declarative, ...)
- Learn how to define and use abstract data types ...)

⇒ Most importantly, recognize that computer science is not programming.

Grading Scheme

- Quizzes (30%)
 - Control your continued learning success (quick feedback)
- Assignments (20%)
 - Learning by solving assignments
 - Test whether you can apply concepts learned
- Mid-term examination (20%)
 - Covers the first half of the course
 - Closed book (and closed computers)
- Final examination (30%)
 - Covers the whole course
 - Closed book (and closed computers)

Teaching and learning strategy

- Quizzes (about 10 minutes) require you to study the material covered in the course continuously
- Homework assignments: Reinforce and apply what is taught in class
- Assignments will be small individual assignments (but may take time to solve)
- Consider forming study groups. It helps to discuss questions and course material in study groups or to explore different directions to solve an assignment. However, solutions must be individual submissions. (Discuss the general problem in a study group, workout the details of the solution yourself.)
- You can audit the course. To earn an audit, you have to take the quizzes and do reasonably well. In addition, there will be a short oral interview about key concepts introduced in the course at the end of the semester.

Organizational aspects and tutorials

- All homework assignments will be linked to the course web page.
- Solutions for assignments will be submitted using the online grader system.
<http://grader.eecs.jacobs-university.de/>
- Feedback and grades will be accessible via the grader system as well.
- Tutorials groups lead by a teaching assistant will be offered to discuss course topics and to raise questions concerning homeworks, quizzes, etc.

Rules of the Game

1 Course Objectives and Grading

2 Rules of the Game

3 Resources

Code of Academic Integrity

- Jacobs University has a “Code of Academic Integrity”
 - this is a document approved by the Jacobs community
 - you have signed it during enrollment
 - it is a law of the university, we take it seriously
- It mandates good behaviours from faculty and students and it penalizes bad ones:
 - honest academic behavior (e.g., no cheating)
 - respect and protect intellectual property of others (e.g., no plagiarism)
 - treat all Jacobs members equally (e.g., no favoritism)
- It protects you and it builds an atmosphere of mutual respect
 - we treat each other as reasonable persons
 - the other's requests and needs are reasonable until proven otherwise
 - if others violate our trust, we are deeply disappointed (may be leading to severe and uncompromising consequences)

Academic Integrity Committee (AIC)

- The Academic Integrity Committee is a joint committee by students and faculty.
- Mandate: to hear and decide on any major or contested allegations, in particular,
 - the AIC decides based on evidence in a timely manner,
 - the AIC makes recommendations that are executed by academic affairs,
 - the AIC tries to keep allegations against faculty anonymous for the student.
- Every member of Jacobs University (faculty, student, ...) can appeal any academic integrity allegations to the AIC.

Cheating

- There is no need to cheat, cheating prevents you from learning.
- Useful collaboration versus cheating:
 - You will be required to hand in your own original code/text/math for all assignments
 - You may discuss your homework assignments with others, but if doing so impairs your ability to write truly original code/text/math, you will be cheating
 - Copying from peers, books or the Internet is plagiarism unless properly attributed
- What happens if we catch you cheating?
 - We will confront you with the allegation (you can explain yourself)
 - If you admit or are silent, we impose a grade sanction and we notify the student records office
 - Repeated infractions to go the AIC for deliberation
- Note: Both active (copying from others) and passive cheating (allowing others to copy) are penalized equally

Deadlines

- Deadlines will be strict (don't bother to ask for extensions)
- Make sure you submit the right document. We grade what was submitted, not what could have been submitted.
- Submit early - avoid last minute changes or software/hardware problems.
- Official excuses by the student records office will extend the deadlines - but not more than the time covered by the excuse.
- A word on medical excuses: Use them when you are ill. Do not use them as a tool to gain more time.
- You want to be taken serious if you are seriously ill. Misuse of excuses can lead to a situation where you are not taken very serious when you deserve to be taken serious.

Culture of Questions, Answers, and Explanations

- Answers to questions require an explanation even if this is not stated explicitly
 - A question like 'Does this algorithm always terminate?' can in principle be answered with 'yes' or 'no'.
 - We expect, however, that an explanation is given why the answer is 'yes' or 'no', even if this is not explicitly stated.
- Answers should be written in your own words
 - Sometimes it is possible to find perfect answers on Wikipedia or Stack Exchange or in good old textbooks.
 - Simply copying the answer of someone else is plagiarism.
 - Copying the answer and providing the reference solves the plagiarism issue but usually does not show that you understood the answer.
 - Hence, we want you to write the answer in your own words.
 - Learning how to write concise and precise answers is an important academic skill.

Culture of Interaction

- I am here to help you learn the material.
- If things are unclear, ask questions.
- If I am going too fast, tell me.
- If I am going too slow, tell me.
- Discussion in class is most welcome - don't be shy.
- Discussion in tutorials is even more welcome - don't be shy.
- If you do not understand something, chances are pretty high your neighbor does not understand either.
- Don't be afraid of asking teaching assistants or myself for help and additional explanations.

1 Course Objectives and Grading

2 Rules of the Game

3 Resources

Study Material and Forums

- There is no required textbook.
- The slides and notes are available on the course web page.
<http://cnds.eecs.jacobs-university.de/courses/ics-2017>
- We encourage you to sign up for the Piazza course forum.
<https://piazza.com/>
- General questions should be asked on the Piazza forum.
 - Faster responses since many people can answer
 - Better responses since people can collaborate on the answer
- For individual questions, send email or come to see me at my office (or talk to me after class or wherever you find me).

Software Tools

- You will need a computer to follow this course.
- Get used to standard software tools:
 - Good and powerful editors such as `emacs` or `vim`
 - Unix-like operating systems such as `Linux`
 - Learn how to use a command interpreter (shell) like `bash` or `zsh`
 - Learn to write structured documents using \LaTeX (great for typesetting math)
 - Learn how to maintain an agenda and TODO items (managing your time)
- Learn how to touch-type (typing without having to look at the keys)

Part: Introduction and Examples

- 4 Computer Science and Algorithms
- 5 Maze Generation Algorithms
- 6 String Search Algorithms
- 7 Complexity, Correctness, Engineering

4 Computer Science and Algorithms

5 Maze Generation Algorithms

6 String Search Algorithms

7 Complexity, Correctness, Engineering

- Computer science is the study of the theory, experimentation, and engineering that form the basis for the design and use of computers
- It is the systematic study of the feasibility, structure, expression, and mechanization of the methodical procedures (or algorithms) that underlie the acquisition, representation, processing, storage, communication of, and access to information
- Computer science is the study of automating algorithmic processes that scale. A computer scientist specializes in the theory of computation and the design of computational systems

Source: [Wikipedia, accessed 2017-06-11]

Definition (algorithm)

In computer science, an *algorithm* is a self-contained sequence of actions to be performed in order to achieve a certain task.

- If you are confronted with a problem, do the following steps:
 - first think about the problem to make sure you fully understand it
 - afterwards try to find an algorithm to solve the problem
 - try to assess the properties of the algorithms (will it handle corner cases correctly? how long will it run? will it always terminate?, ...)
 - consider possible alternatives that may have “better” properties
 - finally, write a program to implement the most suitable algorithm you have selected
- Is the above an algorithm to find algorithms?

Algorithmic Thinking

Algorithmic thinking is a collection of abilities that are essential for constructing and understanding algorithms:

- the ability to analyze given problems
- the ability to specify a problem precisely
- the ability to find the basic actions that are adequate to the given problem
- the ability to construct a correct algorithm using the basic actions
- the ability to think about all possible special and normal cases of a problem
- the ability to assess and improve the efficiency of an algorithm

Maze Generation Algorithms


4 Computer Science and Algorithms

5 Maze Generation Algorithms

6 String Search Algorithms

7 Complexity, Correctness, Engineering

Maze (33 x 11)




Problem Statement

Problem:

- Write a program to generate mazes.
- Every maze should be solvable, i.e., it should have a path from the entrance to the exit.
- We want maze solutions to be unique.
- We want every “room” to be reachable.

Questions:

- How do we approach this problem?
- Are there other properties that make a maze a “good” or a “challenging” maze?




Hacking...



Problem Formalization (1/3)

- Think of a maze as a (two-dimensional) grid of rooms separated by walls.
- Each room can be given a name.
- Initially, every room is surrounded by four walls
- General idea:
 - Randomly knock out walls until we get a good maze.
 - How do we ensure there is a solution?
 - How do we ensure there is a unique solution?
 - How do we ensure every room is reachable?



Problem Formalization (2/3)


Lets try to formalize the problem in mathematical terms:

- We have a set V of rooms.
- We have a set E of pairs (x, y) with $x \in V$ and $y \in V$ of adjacent rooms that have an open wall between them.

In the example, we have

- $V = \{a, b, c, d, e, f, g, h, i, j, k, l, m, n, o, p\}$
- $(a, b) \in E$ and $(g, k) \in E$ and $(a, c) \notin E$ and $(e, f) \notin E$

Abstractly speaking, this is a mathematical structure called a graph consisting of a set of vertices (also called nodes) and a set of edges (also called links).



Why use a mathematical formalization?

- Data structures are typically defined as mathematical structures
- Mathematics can be used to reason about the correctness and efficiency of data structures and algorithms
- Mathematical structures make it easier to *think* — to abstract away from unnecessary details and to avoid “hacking”


Problem Formalization (3/3)

Definition:

- A maze is a graph $G = (V, E)$ with two special nodes, the start node S and the exit node X .


Interpretation:

- Each graph node $x \in V$ represents a room
- An edge $(x, y) \in E$ indicates that rooms x and y are adjacent and there is no wall in between them
- The first special node is the start of the maze
- The second special node is the exit of the maze



Mazes as Graphs (Visualization via Diagrams)

- Graphs are very abstract objects, we need a good, intuitive way of thinking about them.
- We use diagrams, where the nodes are visualized as circles and the edges as lines between them.
- Note that the diagram is a *visualization* of the graph, and not the graph itself.
- A *visualization* is a representation of a structure intended for humans to process visually.



Mazes as Graphs (Good Mazes)


Recall, what is a good maze?

- We want maze solutions to be unique.
- We want every room to be reachable.

Solution:

- The graph must be a tree (a graph with a unique root node and every node except the root node having a unique parent).
- The tree should cover all nodes (we call such a tree a spanning tree).

Since trees have no cycles, we have a unique solution.



Kruskal's Algorithm (1/2)

General approach:

- Randomly add a branch to the tree if it won't create a cycle (i.e., tear down a wall).
- Repeat until a spanning tree has been created.

Questions:

- When adding a branch (edge) (x, y) to the tree, how do we detect that the branch won't create a cycle?
- When adding an edge (x, y) , we want to know if there is already a path from x to y in the tree (if there is one, do not add the edge (x, y)).
- How can we quickly determine whether there is already a path from x to y ?

Kruskal's Algorithm (2/2)

The Union Find Algorithm successively puts nodes into an *equivalence class* if there is a path connecting them. With this idea, we get the following algorithm to construct a spanning tree:

1. Initially, every node is in its own equivalence class and the set of edges is empty.
2. Randomly select a possible edge (x, y) such that x and y are not in the same equivalence class.
3. Add the edge (x, y) to the tree and join the equivalence classes of x and y .
4. Repeat the last two steps if there are still multiple equivalence classes.

Randomized Depth-first Search

Are there other algorithms? Of course there are. Here is a different approach to build a tree rooted at the start node.

1. Make the start node the current node and mark it as visited.
2. While there are unvisited nodes:
 - 2.1 If the current node has any neighbours which have not been visited:
 - 2.1.1 Choose randomly one of the unvisited neighbours
 - 2.1.2 Push the current node to the stack (of nodes)
 - 2.1.3 Remove the wall between the current node and the chosen node
 - 2.1.4 Make the chosen node the current node and mark it as visited
 - 2.2 Else if the stack is not empty:
 - 2.2.1 Pop a node from the stack (of nodes)
 - 2.2.2 Make it the current node

String Search Algorithms

4 Computer Science and Algorithms

5 Maze Generation Algorithms

6 String Search Algorithms

7 Complexity, Correctness, Engineering

Problem Statement

Problem:

- Write a program to find a (relatively short) string in a (possibly long) text.
- This is sometimes called finding a needle in a haystack.

Questions:

- How can we do this efficiently?
- What do we mean with long?
- What exactly is a string and what is text?

Problem Formalization

- Let Σ be a finite set, called an alphabet.
- Let k denote the number of elements in Σ .
- Let Σ^* be the set of all words that can be created out of Σ (Kleene closure of Σ).
- Let $t \in \Sigma^*$ be a (possible long) text and $p \in \Sigma^*$ be a (typically short) pattern.
- Let n denote the length of t and m denote the length of p .
- We assume that $n \gg m$.
- Find the first occurrence of p in t .

Naive String Search

- Check whether the pattern matches at each text position (going left to right).
- Lowercase characters indicate comparisons that were skipped.
- Example: $t = \text{FINDANEEDLEINAHAYSTACK}$, $p = \text{NEEDLE}$

F I N D A N E E D L E I N A H A Y S T A C K

N e e d l e

N e e d l e

N E e d l e

N e e d l e

N e e d l e

N E E D L E

Naive String Search Performance

- How “fast” is naive string search?
- Idea: Lets try to count the number of comparisons.
- Problem: The number of comparisons depends on the strings.
- Idea: Consider the worst case possible.
- What is the worst case possible?
 - Consider a haystack of length n using only a single symbol of the alphabet (e.g., “aaaaaaaaaa” with $n = 10$).
 - Consider a needle of length m which consists of $m - 1$ times the same symbol followed by a single symbol that is different (e.g., “aax” with $m = 3$).
- With $n \gg m$, the number of comparisons needed will be roughly $n \cdot m$.

Boyer-Moore: Bad character rule (1/2)

- Idea: Lets compare the pattern right to left instead left to right. If there is a mismatch, try to move the pattern as much as possible to the right.
- Bad character rule: Upon mismatch, move the pattern to the right until there is a match at the current position or until the pattern has moved past the current position.
- Example: $t = \text{FINDANEEDLEINAHAYSTACK}$, $p = \text{NEED}$

F I N D A N E E D L E I N A H A Y S T A C K	skip
n e E D	1
n e e D	2
N E E D	

Boyer-Moore: Bad character rule (2/2)

- Example: $t = \text{FINDANEEDLEINAHAYSTACK}$, $p = \text{HAY}$

F I N D A N E E D L E I N A H A Y S T A C K	skip
h a Y	2
h a Y	2
h a Y	2
h a Y	2
h a Y	1
H A Y	

- How do we decide efficiently how far we can move the pattern to the right?

Boyer-Moore: Good suffix rule (1/3)

- Idea: If we already matched a suffix and the suffix appears again in the pattern, skip the alignment such that we keep the good suffix.
- Good suffix rule: Let s be the suffix already matched in the inner loop. If there is a mismatch, skip alignments until (i) there is another match of the suffix, or (ii) a prefix of p matches a suffix of s or (iii) if there is no such suffix, skip until the end of the pattern.
- Example: $t = \text{FINDANEEDLEINAHAYSTACK}$, $p = \text{NEEDUNEED}$

F I N D A N E E D L E I N A H A Y S T A C K	skip
n e e d U N E E D	4
n e e d u n e e D	

Boyer-Moore: Good suffix rule (2/3)

- Example: $t = \text{FINDANEEDLEINAHAYSTACK}$, $p = \text{EDISUNEEED}$

F I N D A N E E D L E I N A H A Y S T A C K skip

e d i s U N E E D 6

 e d i s u n e e D

Boyer-Moore: Good suffix rule (3/3)

- Example: $t = \text{FINDANEEDLEINAHAYSTACK}$, $p = \text{FOODINEED}$

F I N D A N E E D L E I N A H A Y S T A C K skip

f o o d I N E E D 8

f o o d i n e e D

- How do we decide efficiently how far we can move the pattern to the right?

Boyer-Moore Rules Combined

- The Boyer-Moore algorithm combines the bad character rule and the good suffix rule. (Note that both rules can also be used alone.)
- If a mismatch is found,
 - calculate the skip s_b by the bad character rule
 - calculate the skip s_g by the good suffix ruleand then skip by $s = \max(s_b, s_g)$.
- The Boyer-Moore algorithm often does the substring search in sub-linear time.
- However, it does not perform better than naive search in the worst case if the pattern does occur in the text.
- An optimization by Gali results in linear runtime across all cases.

Complexity, Correctness, Engineering

- 4 Computer Science and Algorithms
- 5 Maze Generation Algorithms
- 6 String Search Algorithms
- 7 Complexity, Correctness, Engineering**

Complexity of Algorithms

- Questions:
 - Which maze generation algorithm is faster?
 - Is there a fastest maze generation algorithm?
 - What happens if we consider mazes of different sizes or dimensions?
 - Instead of measuring execution time (which depends on the speed of the computer hardware), can we have a more neutral notion of “fast”?
- Computer science is about analyzing the complexity of algorithms.
- Complexity is an abstract measure of computational effort (time complexity) and memory usage (space complexity).

Performance and Scaling

- Suppose we have three algorithms to choose from. For example, consider algorithms to detect cycles in a graph of n nodes.
- With $n = 50$, the exponential algorithm has an execution time of more than 35 years.
- For $n \geq 1000$, the exponential algorithm gives us execution times that are longer than the age of the universe!

size n	linear $100n \mu\text{s}$	quadratic $7n^2 \mu\text{s}$	exponential $2^n \mu\text{s}$
1	100 μs	7 μs	2 μs
5	500 μs	175 μs	32 μs
10	1 ms	700 μs	1024 μs
50	5 ms	17.5 ms	13 031.25 d
100	10 ms	70 ms	
1000	100 ms	7 s	
10 000	1 s	700 s	
100 000	10 s	70 000 s	

Correctness of Algorithms and Programs

- Questions:
 - Is our algorithm correct?
 - Is our algorithm a total function or a partial function?
 - Is our implementation of the algorithm (our program) correct?
 - What do we mean by “correct”?
 - Will our algorithm or program terminate?
- Computer science is about techniques for proving correctness of programs.
- In situations where correctness proofs are not feasible, computer sciences is about engineering practices that help to avoid or detect errors.

Partial Correctness and Total Correctness

Definition (partial correctness)

An algorithm starting in a state that satisfies a precondition P is *partially correct with respect to P and Q* if results produced by the algorithm satisfy the postcondition Q . Partial correctness does not require that always a result is produced, i.e., the algorithm may not always terminate.

Definition (total correctness)

An algorithm is *totally correct with respect to P and Q* if it is partially correct with respect to P and Q and it always terminates.

Deterministic Algorithms

Definition (deterministic algorithm)

A *deterministic algorithm* is an algorithm which, given a particular input, will always produce the same output, with the underlying machine always passing through the same sequence of states.

- Some factors that make an algorithm non-deterministic:
 - external state
 - user input
 - timers
 - random values
 - hardware errors

Randomized Algorithms

Definition (randomized algorithm)

A *randomized algorithm* is an algorithm that employs a degree of randomness as part of its logic.

- A randomized algorithm uses randomness in order to produce its result; it uses randomness as part of the logic of the algorithm.
- A perfect source of randomness is not trivial to obtain on digital computers.
- Random number generators often use algorithms to produce so called pseudo random numbers, sequences of numbers that “look” random but that are not really random (since they are calculated using a deterministic algorithm).

- Questions:
 - Can we identify building blocks (data structures, generic algorithms, design pattern) that we can reuse?
 - Can we implement algorithms in such a way that the program code is easy to read and understand?
 - Can we implement algorithms in such a way that we can easily adapt them to different requirements?
- Computer science is about modular designs that are both easier to get right and easier to understand. Finding good software designs often takes time and effort.
- Software engineering is about applying structured approaches to the design, development, maintenance, testing, and evaluation of software.

Part: Discrete Mathematics

8 Propositions, Axioms, Theorems, Proofs

9 Sets

10 Relations

11 Functions

Propositions, Axioms, Theorems, Proofs

8 Propositions, Axioms, Theorems, Proofs

9 Sets

10 Relations

11 Functions

Propositions

Definition (proposition)

A *proposition* is a statement that is either true or false.

Examples:

- $1 + 1 = 1$ (false proposition)
- The sum of the integer numbers $1, \dots, n$ is equal to $\frac{1}{2}n(n+1)$. (true proposition)
- “In three years I will have obtained a CS degree.” (not a proposition)

- A predicate is a statement that may be true or false depending on the values of its variables. It can be thought of as a function that returns a value that is either true or false. Variables appearing in a predicate are quantified:
 - A predicate is true for all values of a given set of values.
 - A predicate is true for at least one value of a given set of values.
(There exists a value such that the predicate is true.)
- There may be multiple quantifiers and they may be combined (but note that the order of the quantifiers matters).
- Example: (Goldbach's conjecture) For every even integer n greater than 2, there exists primes p and q such that $n = p + q$.

Definition (axiom)

An *axiom* is a proposition that is taken to be true.

Definition (Peano axioms for natural numbers)

- P1 0 is a natural number.
- P2 Every natural number has a successor.
- P3 0 is not the successor of any natural number.
- P4 If the successor of x equals the successor of y , then x equals y .
- P5 If a statement is true for the natural number 0, and if the truth of that statement for a natural number implies its truth for the successor of that number, then the statement is true for every natural number.

Definition (theorem, lemma, corollary)

An important true proposition is called a *theorem*.

A *lemma* is a preliminary proposition useful for proving later propositions and a *corollary* is a proposition that follows in just a few logical steps from a theorem.

- There is no clear boundary between what is a theorem, a lemma, or a corollary.
- A proposition for which no proof has been found yet and which is believed to be true is called a *conjecture*.

Mathematical Notation

Notation	Explanation
$P \wedge Q$	logical and of propositions P and Q
$P \vee Q$	logical or of propositions P and Q
$\neg P$	negation of proposition P
$\forall x \in S. P$	the predicate P holds for all x in the set S
$\exists x \in S. P$	there exists an x in the set S such that the predicate P holds
$P \Rightarrow Q$	the statement P implies statement Q
$P \Leftrightarrow Q$	the statement P holds if and only if (iff) Q holds

Greek Letters

α	A	alpha	β	B	beta	γ	Γ	gamma
δ	Δ	delta	ε	E	epsilon	ζ	Z	zeta
η	H	eta	θ	Θ	theta	ι	I	iota
κ	K	kappa	λ	Λ	lambda	μ	M	mu
ν	N	nu	ξ	Ξ	xi	\omicron	O	omikron
π	Π	pi	ρ	ρ	rho	σ	Σ	sigma
τ	T	tau	υ	Υ	upsilon	φ	Φ	phi
χ	X	chi	ψ	Ψ	psi	ω	Ω	omega

Definition (mathematical proof)

A *mathematical proof* of a proposition is a chain of logical deductions from a base set of axioms (or other previously proven propositions) that concludes with the proposition in question.

- Informally, a proof is a method of establishing truth. There are very different ways to establish truth. In computer science, we usually adopt the mathematical notion of a proof.
- There are a certain number of templates for constructing proofs. It is good style to indicate at the beginning of the proof which template is used.

Proofs Hints

- Proofs often start with scratchwork that can be disorganized, have strange diagrams, obscene words, whatever. But the final proof should be clear and concise.
- Proofs usually begin with the word “Proof” and they end with a delimiter such as \square .
- Make it easy to understand your proof. A good proof has a clear structure and it is concise.
- Introduce notation carefully. Good notation can make a proof easy to follow (and bad notation can achieve the opposite effect).
- Revise your proof and simplify it. A good proof has been written multiple times.

Proof an Implication by Derivation

- An implication is a proposition of the form “If P , then Q ”, or $P \Rightarrow Q$.
- One way to proof such an implication is by a derivation where you start with P and stepwise derive Q from it.
- In each step, you apply theorems (or lemmas or corollaries) that have already been proven to be true.
- Template:
Assume P . Then, ... Therefore ... [...] This finally leads to Q . \square

Proof an Implication by its Contrapositive

- An implication is a proposition of the form “If P , then Q ”, or $P \Rightarrow Q$.
- Such an implication is logically equivalent to its *contrapositive*, $\neg Q \Rightarrow \neg P$.
- Proving the contrapositive is sometimes easier than proving the original statement.
- Template:

Proof. We prove the contrapositive, if $\neg Q$, then $\neg P$. We assume $\neg Q$. Then, ... Therefore ... [...] This finally leads to $\neg P$. \square

Proof an “if and only if” by two Implications

- A statement of the form “ P if and only if Q ” is equivalent to the two statements “ P implies Q ” and “ Q implies P ”.
- Split your proof into two parts, the first part proving $P \Rightarrow Q$ and the second part proving $Q \Rightarrow P$.

- Template:

Proof. We prove P implies Q and vice-versa.

First, we show P implies Q . Assume P . Then, ... Therefore ... [...] This finally leads to Q .

Now we show Q implies P . Assume Q . Then, Therefore ... [...] This finally leads to P . \square

Proof an “if and only if” by a Chain of “if and only if” s

- A statement of the form “ P if and only if Q ” can be shown to hold by constructing a chain of “if and only if” equivalence implications.
- Constructing this kind of proof is often harder then proving two implications, but the result can be short and elegant.

- Template:

Proof. We construct a proof by a chain of if-and-only-if implications.

Prove P is equivalent to a P' which is equivalent to $[\dots]$ which is equivalent to Q .



Breaking a Proof into Cases

- It is sometimes useful to break a complicated statement P into several cases that are proven separately.
- Different proof techniques may be used for the different cases.
- It is necessary to ensure that the cases cover the complete statement P .
- Template:

Proof. We prove P by considering the cases c_1, \dots, c_N .

Case 1: Suppose c_1 . Proof of P for c_1 .

...

Case N : Suppose c_N . Proof of P for c_N .

Since P holds for all cases c_1, \dots, c_N hold, the statement P holds. \square

Proof by Contradiction

- A proof by contradiction for a statement P shows that if the statement were false, then some false fact would be true.
- Starting from $\neg P$, a series of derivations is used to arrive at a statement that contradicts something that has already been shown to be true or which is an axiom.
- Template:

Proof. We prove P by contradiction.

Assume $\neg P$ is true. Then ... Therefore ... [...] This is a contradiction. Thus, P must be true. \square

Proof by Induction

- If we have to prove a statement P on nonnegative integers (or more generally an inductively defined infinite set), we can use the induction principle.
- We first prove that P is true for the “lowest” element in the set (the base case).
- Next we prove that if P holds for a nonnegative integer n , then the statement P holds for $n + 1$ (induction step).
- Since we can apply the induction step m times, starting with the base, we have shown that P is true for arbitrary nonnegative integers m .

- Template:

Proof. We prove P by induction.

Base case: We show that $P(0)$ is true. [...]

Induction step: Assume $P(n)$ is true. Then, ... This proves that $P(n + 1)$ holds.

8 Propositions, Axioms, Theorems, Proofs

9 Sets

10 Relations

11 Functions

- Informally, a *set* is a well-defined collection of distinct objects. The elements of the collection can be anything we like the set to contain, including other sets.
- In modern math, sets are defined using axiomatic set theory, but for us the informal definition above is sufficient.
- Sets can be defined by
 - listing all elements in curly braces, e.g., $\{a, b, c\}$,
 - describing all objects using a predicate P , e.g., $\{x | x \geq 0 \wedge x < 2^8\}$,
 - stating element-hood using some other statements.
- A set has no order of the elements and every element appears only once.
- The two notations $\{a, b, c\}$ and $\{b, a, a, c\}$ are different *representations* of the same set.

Basic Relations between Sets

Definition (basic relations between sets)

Lets A and B be two sets. We define the following relations between sets:

1. $(A \equiv B) :\Leftrightarrow (\forall x. x \in A \Leftrightarrow x \in B)$ (set equality)
2. $(A \subseteq B) :\Leftrightarrow (\forall x. x \in A \Rightarrow x \in B)$ (subset)
3. $(A \subset B) :\Leftrightarrow (A \subseteq B) \wedge (A \neq B)$ (proper subset)
4. $(A \supseteq B) :\Leftrightarrow (\forall x. x \in B \Rightarrow x \in A)$ (superset)
5. $(A \supset B) :\Leftrightarrow (A \supseteq B) \wedge (A \neq B)$ (proper superset)

- Obviously:
 - $(A \subseteq B) \wedge (B \subseteq A) \Rightarrow (A \equiv B)$
 - $(A \subseteq B) \Leftrightarrow (B \supseteq A)$

Operations on Sets 1/2

Definition (set union)

The *union* of two sets A and B is defined as $A \cup B = \{x | x \in A \vee x \in B\}$.

Definition (set intersection)

The *intersection* of two sets A and B is defined as $A \cap B = \{x | x \in A \wedge x \in B\}$.

Definition (set difference)

The *difference* of two sets A and B is defined as $A \setminus B = \{x | x \in A \wedge x \notin B\}$.

Operations on Sets 2/2

Definition (power set)

The *power set* $\mathcal{P}(A)$ of a set A is the set of all subsets of S , including the empty set and S itself. Formally, $\mathcal{P}(A) = \{S \mid S \subseteq A\}$.

Definition (cartesian product)

The *cartesian product* of the sets X_1, \dots, X_n is defined as $X_1 \times \dots \times X_n = \{(x_1, \dots, x_n) \mid \forall i. 1 \leq i \leq n \Rightarrow x_i \in X_i\}$.

Cardinality of Sets

Definition (cardinality)

If A is a finite set, the *cardinality* of A , written as $|A|$, is the number of elements in A .

Definition (countably infinite)

A set A is *countably infinite* if and only if there is a bijective function $f : A \mapsto \mathbb{N}$.

Definition (countable)

A set A is *countable* if and only if it is finite or countably infinite.

8 Propositions, Axioms, Theorems, Proofs

9 Sets

10 Relations

11 Functions

Definition (relation)

A *relation* R over the sets X_1, \dots, X_k is a subset of their Cartesian product, written $R \subseteq X_1 \times \dots \times X_k$.

- Relations are classified according to the number of sets in the defining Cartesian product:
 - A unary relation is defined over a single set X
 - A binary relation is defined over $X_1 \times X_2$
 - A ternary relation is defined over $X_1 \times X_2 \times X_3$
 - A k -ary relation is defined over $X_1 \times \dots \times X_k$

Binary Relations

Definition (binary relation)

A *binary relation* $R \subseteq A \times B$ consists of a set A , called the *domain* of R , a set B , called the *codomain* of R , and a subset of $A \times B$ called the *graph* of R .

Definition (inverse of a binary relation)

The *inverse* of a binary relation $R \subseteq A \times B$ is the relation $R^{-1} \subseteq B \times A$ defined by the rule

$$b R^{-1} a \Leftrightarrow a R b.$$

- For $a \in A$ and $b \in B$, we often write $a R b$ to indicate that $(a, b) \in R$.
- The notation $a R b$ is called *infix notation* while the notation $R(a, b)$ is called the *prefix notation*. For binary relations, we commonly use the infix notation.

Image and Range of Binary Relations

Definition (image of a binary relation)

The *image* of a binary relation $R \subseteq A \times B$, is the set of elements of the codomain B of R that are related to some element in A .

Definition (range of a binary relation)

The *range* of a binary relation $R \subseteq A \times B$ is the set of elements of the domain A of R that relate to at least one element in B .

Properties of Binary Relations (Endorelations)

Definition

A relation $R \subseteq A \times A$ is called

- *reflexive* iff $\forall a \in A. (a, a) \in R$
- *irreflexive* iff $\forall a \in A. (a, a) \notin R$
- *symmetric* iff $\forall a, b \in A. (a, b) \in R \Rightarrow (b, a) \in R$
- *asymmetric* iff $\forall a, b \in A. (a, b) \in R \Rightarrow (b, a) \notin R$
- *antisymmetric* iff $\forall a, b \in A. ((a, b) \in R \wedge (b, a) \in R) \Rightarrow a = b$
- *transitive* iff $\forall a, b, c \in A. ((a, b) \in R \wedge (b, c) \in R) \Rightarrow (a, c) \in R$
- *total* iff $\forall a, b \in A. (a, b) \in R \vee (b, a) \in R$
- *equivalence relation* iff R is reflexive, symmetric, and transitive.

Partial and Strict Order

Definition (partial order and strict partial order)

A relation $R \subseteq A \times A$ is called a *partial order* on A if and only if R is reflexive, antisymmetric, and transitive on A . The relation R is called a *strict partial order* on A if and only if it is irreflexive, asymmetric and transitive on A .

Definition (linear order)

A partial order R is called a *linear order* on A if and only if all elements in A are comparable, i.e., the partial order is total.

- A symbol commonly used for strict partial orders is \prec and a symbol commonly used for non-strict partial orders is \preceq .

8 Propositions, Axioms, Theorems, Proofs

9 Sets

10 Relations

11 Functions

Definition (partial function)

A relation $f \subseteq X \times Y$ is called a *partial function* if and only if for all $x \in X$ there is at most one $y \in Y$ with $(x, y) \in f$. We call a partial function f undefined at $x \in X$ if and only if $(x, y) \notin f$ for all $y \in Y$.

Definition (total function)

A partial function $f \subseteq X \times Y$ is called a total function if f is a total relation.

Function Properties

Definition (injective function)

A function $f : X \mapsto Y$ is called *injective* if every element of the codomain Y is mapped to by at most one element of the domain X : $\forall x, y \in X. f(x) = f(y) \Rightarrow x = y$

Definition (surjective function)

A function $f : X \mapsto Y$ is called *surjective* if every element of the codomain Y is mapped to by *at least one* element of the domain X : $\forall y \in Y. \exists x \in X. f(x) = y$

Definition (bijective function)

A function $f : X \mapsto Y$ is called *bijective* if every element of the codomain Y is mapped to by exactly one element of the domain X . (That is, the function is both injective and surjective.)

Operations on Functions

Definition (function composition)

Given two functions $f : A \mapsto B$ and $g : B \mapsto C$, the *composition* of g with f is defined as the function $g \circ f : A \mapsto C$ with $(g \circ f)(x) = g(f(x))$.

Definition (function restriction)

Let f be a function $f : A \mapsto B$ and $C \subseteq A$. Then we call the function $f|_C = \{(c, b) \in f \mid c \in C\}$ the restriction of f to C .

Lambda Notation of Functions

- It is often not necessary to give a function a name.
- A function definition of the form $\{(x, y) \in X \times Y \mid y = E\}$, where E is an expression (usually involving x), can be written in a shorter lambda notation as $\lambda x \in X. E$.
- Examples:
 - $\lambda n \in \mathbb{N}. n$ (identity function for natural numbers)
 - $\lambda x \in \mathbb{N}. x^2$ ($f(x) = x^2$)
 - $\lambda(x, y) \in \mathbb{N} \times \mathbb{N}. x + y$ (addition of natural numbers)
- Lambda calculus is a formal system in mathematical logic for expressing computation based on function abstraction and application using variable binding and substitution.

Currying

- Lambda calculus uses only functions that take a single argument. This is possible since lambda calculus allows functions as arguments and results.
- A function that takes two arguments can be converted into a function that takes the first argument as input and which returns a function that takes the second argument as input.
- This method of converting function with multiple arguments into a sequence of functions with a single argument is called currying.
- The term currying is a reference to the logician Haskell Curry.

Part: Number Systems, Units, Characters, Date and Time

12 Natural Numbers

13 Integer Numbers

14 Rational and Real Numbers

15 Floating Point Numbers

16 International System of Units

17 Characters and Strings

18 Date and Time

Numbers can be confusing...

- There are only 10 people in the world: Those who understand binary and those who don't.
- Q: How easy is it to count in binary?
A: Its as easy as 01 10 11.
- A Roman walks into the bar, holds up two fingers, and says, "Five beers, please."
- Q: Why do mathematicians confuse Halloween and Christmas?
A: Because $31 \text{ Oct} = 25 \text{ Dec}$.

Number Systems in Mathematics

- Numbers can be classified into sets, called number systems, such as the natural numbers, the integer numbers, or the real numbers.

Symbol	Name	Description
\mathbb{N}	Natural	$0, 1, 2, 3, 4, \dots$
\mathbb{Z}	Integer	$\dots, -4, -3, -2, -1, 0, 1, 2, 3, 4, \dots$
\mathbb{Q}	Rational	$\frac{a}{b}$ where $a \in \mathbb{Z}$ and $b \in \mathbb{Z}$ and $b \neq 0$
\mathbb{R}	Real	The limit of a convergent sequence of rational numbers
\mathbb{C}	Complex	$a + bi$ where $a \in \mathbb{R}$ and $b \in \mathbb{R}$ and $i = \sqrt{-1}$

- Numbers should be distinguished from numerals, the symbols used to represent numbers. A single number can have many different representations.

Natural Numbers

12 Natural Numbers

13 Integer Numbers

14 Rational and Real Numbers

15 Floating Point Numbers

16 International System of Units

17 Characters and Strings

18 Date and Time

Numeral Systems for Natural Numbers

- Natural numbers can be represented according to different bases. We commonly use decimal number (base 10) representations in everyday life.
- In computer science, we also frequently use binary (base 2), octal (base 8), and hexadecimal (base 16) number representations.
- In general, natural numbers represented in the base b system are of the form:

$$(a_n a_{n-1} \cdots a_1 a_0)_b = \sum_{k=0}^n a_k b^k$$

hex	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	10	11	12
dec	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
oct	0	1	2	3	4	5	6	7	10	11	12	13	14	15	16	17	20	21	22
bin	0	1	10	11	100	101	110	111	1000	1001	1010	1011	1100	1101	1110	1111	10000	10001	10010

Natural Numbers Literals

- Computer scientists often use special prefix conventions to write natural number literals in a way that indicates the base:

prefix	example	meaning	description
	42	42_{10}	decimal number
0x	0x42	$42_{16} = 66_{10}$	hexadecimal number
0	042	$42_8 = 34_{10}$	octal number
0b	0b1000010	$1000010_2 = 42_{10}$	binary number

- Beware that 42 and 042 may not represent the same number!

Natural Numbers with Fixed Precision

- Computer systems often work internally with finite subsets of natural numbers.
- The number of bits used for the binary representation defines the size of the subset.

bits	name	range (decimal)	range (hexadecimal)
4	nibble	0-15	0x0-0xf
8	byte, octet, uint8	0-255	0x0-0xff
16	uint16	0-65 535	0x0-0xffff
32	uint32	0-4 294 967 295	0x0-0xffffffff
64	uint64	0-18 446 744 073 709 551 615	0x0-0xffffffffffffffff

- Using (almost) arbitrary precision numbers is possible but usually slower.

Integer Numbers

12 Natural Numbers

13 Integer Numbers

14 Rational and Real Numbers

15 Floating Point Numbers

16 International System of Units

17 Characters and Strings

18 Date and Time

Integer Numbers

- Integer numbers can be negative but surprisingly there are not “more” of them (even though integer numbers range from $-\infty$ to $+\infty$ while natural numbers only range from 0 to $+\infty$).
- This can be seen by writing integer numbers in the order 0, 1, -1, 2, -2, \dots , i.e., by defining a bijective function $f : \mathbb{Z} \rightarrow \mathbb{N}$ (and the inverse function $f^{-1} : \mathbb{N} \rightarrow \mathbb{Z}$):

$$f(x) = \begin{cases} 2x & \text{if } x \geq 0 \\ -2x - 1 & \text{if } x < 0 \end{cases} \quad f^{-1}(x) = \begin{cases} \frac{x}{2} & \text{if } x \text{ is even} \\ \frac{-(x+1)}{2} & \text{if } x \text{ is odd} \end{cases}$$

- So we could (in principle) represent integer numbers by implementing a bijection to natural numbers. But there are more efficient ways to implement integer numbers if we assume that we use a fixed precision anyway.

One's Complement Fixed Integer Numbers ($b-1$ complement)

- We have a fixed number space with n digits and base b to represent integer numbers, that is, we can distinguish at most b^n different integers.
- Lets represent the numbers $0 \dots b^{n-1}$ in the usual way.
- To represent negative numbers, we invert the absolute value $(a_n a_{n-1} \dots a_1 a_0)_b$ by calculating $(a'_n a'_{n-1} \dots a'_1 a'_0)_b$ with $a'_i = (b - 1) - a_i$.
- Example: $b = 2, n = 4 : 5_{10} = 0101, -5_{10} = 1010$

bin:	0000	0001	0010	0011	0100	0101	0110	0111	1000	1001	1010	1011	1100	1101	1110	1111
dec:	0	1	2	3	4	5	6	7	-7	-6	-5	-4	-3	-2	-1	-0

- Note that this gives us $+0$ and -0 , i.e., we only represent $b^n - 1$ different integers.
- Negative binary numbers always have the most significant bit set to 1.

Two's Complement Fixed Integer Numbers (b complement)

- Like before, we assume a fixed number space with n digits and a base b to represent integer numbers, that is, we can distinguish at most b^n different integers.
- Lets again represent the numbers $0 \dots b^{n-1}$ in the usual way.
- To represent negative numbers, we invert the absolute value $(a_n a_{n-1} \dots a_1 a_0)_b$ by calculating $(a'_n a'_{n-1} \dots a'_1 a'_0)_b$ with $a'_i = (b - 1) - a_i$ and adding 1 to it.
- Example: $b = 2, n = 4 : 5_{10} = 0101, -5_{10} = 1011$

bin:	0000	0001	0010	0011	0100	0101	0110	0111	1000	1001	1010	1011	1100	1101	1110	1111
dec:	0	1	2	3	4	5	6	7	-8	-7	-6	-5	-4	-3	-2	-1

- This representation simplifies the implementation of arithmetic operations.
- Negative binary numbers always have the most significant bit set to 1.

Two's Complement Fixed Integer Number Ranges

- Most computers these days use the two's complement internally.
- The number of bits available defines the ranges we can use.

bits	name	range (decimal)
8	int8	−128 to 127
16	int16	−32 768 to 32 767
32	int32	−2 147 483 648 to 2 147 483 647
64	int64	−9 223 372 036 854 775 808 to 9 223 372 036 854 775 807

- Be careful if your arithmetic expressions overflows/underflows the range!

Rational and Real Numbers

12 Natural Numbers

13 Integer Numbers

14 Rational and Real Numbers

15 Floating Point Numbers

16 International System of Units

17 Characters and Strings

18 Date and Time

Rational Numbers

- Computer systems usually do not natively represent rational numbers, i.e., they cannot compute with rational numbers at the hardware level.
- Software can, of course, implement rational number data types by representing the numerator and the denominator as integer numbers internally and keeping them in the reduced form.
- Example using Haskell (execution prints 5 % 6):

```
import Data.Ratio
print $ 1%2 + 1%3
```

Real Numbers

- Computer systems usually do not natively represent real numbers, i.e., they cannot compute with real numbers at the hardware level.
- The primary reason is that real numbers like the result of $\frac{1}{7}$ or numbers like π have by definition not a finite representation.
- So the best we can do is to have a finite approximation. . .
- Since all we have are approximations of real numbers, we *always* make rounding errors if we use these approximations. If we are not very careful, these rounding errors can *accumulate* badly.
- The notion of *numeric stability* can be used to classify algorithms according how they propagate rounding errors.

Floating Point Numbers

12 Natural Numbers

13 Integer Numbers

14 Rational and Real Numbers

15 Floating Point Numbers

16 International System of Units

17 Characters and Strings

18 Date and Time

Floating Point Numbers

- Floating point numbers are useful in situations where a large range of numbers must be represented with fixed size storage for the numbers.
- The general notation of a floating point number f is

$$f = s \times d_0.d_1d_2 \dots d_{p-1} \times b^e$$

where b is the basis, e is the exponent, d_0, d_1, \dots, d_{p-1} are digits of the mantissa with $d_i \in \{0, \dots, b-1\}$ for $i \in \{0, \dots, p-1\}$, $s \in \{1, -1\}$ is the sign, and p is the precision.

Floating Point Number Normalization

- Floating point numbers are usually normalized such that d_0 is in the range $\{1, \dots, b - 1\}$, except when the number is zero.
- Normalization must be checked and restored after each arithmetic operation since the operation may denormalize the number.
- When using the base $b = 2$, normalization implies that the first digit d_0 is always 1. Hence, it is not necessary to store d_0 and instead the mantissa can be extended by one additional bit.
- Floating point numbers are at best an approximation of a real number due to the limited precision.
- Calculations involving floating point numbers usually do not lead to precise results since rounding must be used to match the result into the floating point format.

IEEE 754 Floating Point Formats

Item	Single precision	Double precision	Quad precision
sign	1 bit	1 bit	1 bit
exponent	8 bit	11 bit	15 bit
mantissa	23 bit	52 bit	112 bit
total size	32 bit	64 bit	128 bit
decimal digits	≈ 7.2	≈ 15.9	≈ 34.0

- IEEE 754 is a standard for floating point numbers that is widely implemented today.
- IEEE 754 floating point numbers use the base $b = 2$ and as a consequence numbers such as 1×10^{-1} cannot be represented precisely.

IEEE 754 Exceptions and Special Values

- The standard defines five exceptions, some of them lead to special values:
 1. Invalid operation: returns not a number (nan)
 2. Division by zero: returns \pm infinity (inf)
 3. Overflow: returns \pm infinity (inf)
 4. Underflow: depends on the operating mode
 5. Inexact: returns rounded result by default
- Note that computations may continue if you hit a special value like nan or inf.
- Hence, it is important to check whether a calculation resulted in a value at all.

Floating Point Surprises

- Any floating point computation should be treated with the utmost suspicion unless you can argue how accurate it is. [Alan Mycroft, Cambridge]
- Floating point arithmetic almost always involves rounding errors and these errors can badly aggregate.
- It is possible to “loose” the reasonably precise digits and to continue calculation with the remaining rather imprecise digits.
- Comparisons to floating point constants may not be “exact” and as a consequence loops may not end where they are expected to end.

International System of Units

12 Natural Numbers

13 Integer Numbers

14 Rational and Real Numbers

15 Floating Point Numbers

16 International System of Units

17 Characters and Strings

18 Date and Time

Importance of Units and Unit Prefixes

- Most numbers we encounter in practice have associated units. It is important to be very clear about the units used.
 - NASA lost a Mars climate orbiter (worth \$125 million) in 1999 due to a unit conversion error.
 - An Air Canada plane ran out of fuel in the middle of a flight in 1983 due to a fuel calculation error while switching to the metric system.
 - There is an International System of Units (SI Units) to help you...
- ▶ Always be clear about units.
- ▶ And always be clear about the unit prefixes.

SI Base Units

Unit	Symbol	Description
metre	m	The distance travelled by light in a vacuum in a certain fraction of a second.
kilogram	kg	The mass of the international prototype kilogram.
second	s	The duration of a number of periods of the radiation of the caesium-133 atom.
ampere	A	The constant electric current which would produce between two conductors a certain force.
kelvin	K	A fraction of the thermodynamic temperature of the triple point of water.
mole	mol	The amount of substance of a system which contains atoms corresponding to a certain mass of carbon-12.
candela	cd	The luminous intensity of a source that emits monochromatic radiation.

SI Derived Units

- Many important units can be derived from the base units. Some have special names, others are simply defined by their units. Some examples:

Name	Symbol	Definition	Description
herz	Hz	s^{-1}	frequency
newton	N	$kg\ m\ s^{-1}$	force
watt	W	$kg\ m^2\ s^{-3}$	power
volt	V	$kg\ m^2\ s^{-3}\ A^{-1}$	voltage
ohm	Ω	$kg\ m^2\ s^{-2}\ A^{-1}$	resistance
velocity		$m\ s^{-1}$	speed

Metric Prefixes (International System of Units)

Name	Symbol	Base 10	Base 1000	Value
kilo	k	10^3	1000^1	1000
mega	M	10^6	1000^2	1 000 000
giga	G	10^9	1000^3	1 000 000 000
tera	T	10^{12}	1000^4	1 000 000 000 000
peta	P	10^{15}	1000^5	1 000 000 000 000 000
exa	E	10^{18}	1000^6	1 000 000 000 000 000 000
zetta	Ζ	10^{21}	1000^7	1 000 000 000 000 000 000 000
yotta	Υ	10^{24}	1000^8	1 000 000 000 000 000 000 000 000

Metric Prefixes (International System of Units)

Name	Symbol	Base 10	Base 1000	Value
milli	m	10^{-3}	1000^{-1}	0.001
micro	μ	10^{-6}	1000^{-2}	0.000 001
nano	n	10^{-9}	1000^{-3}	0.000 000 001
pico	p	10^{-12}	1000^{-4}	0.000 000 000 001
femto	f	10^{-15}	1000^{-5}	0.000 000 000 000 001
atto	a	10^{-18}	1000^{-6}	0.000 000 000 000 000 001
zepto	z	10^{-21}	1000^{-7}	0.000 000 000 000 000 000 001
yocto	y	10^{-24}	1000^{-8}	0.000 000 000 000 000 000 000 001

Binary Prefixes

Name	Symbol	Base 2	Base 1024	Value
kibi	Ki	2^{10}	1024^1	1024
mebi	Mi	2^{20}	1024^2	1 048 576
gibi	Gi	2^{30}	1024^3	1 073 741 824
tebi	Ti	2^{40}	1024^4	1 099 511 627 776
pebi	Pi	2^{50}	1024^5	1 125 899 906 842 624
exbi	Ei	2^{60}	1024^6	1 152 921 504 606 846 976
zebi	Zi	2^{70}	1024^7	1 180 591 620 717 411 303 424
yobi	Yi	2^{80}	1024^8	1 208 925 819 614 629 174 706 176

Characters and Strings

12 Natural Numbers

13 Integer Numbers

14 Rational and Real Numbers

15 Floating Point Numbers

16 International System of Units

17 Characters and Strings

18 Date and Time

Characters and Character Encoding

- A *character* is a unit of information that roughly corresponds to a grapheme, grapheme-like unit, or symbol, such as in an alphabet or syllabary in the written form of a natural language.
- Examples of characters include letters, numerical digits, common punctuation marks, and whitespace.
- Characters also includes control characters, which do not correspond to symbols in a particular natural language, but rather to other bits of information used to control information flow or presentation.
- A *character encoding* is used to represent a set of characters by some kind of encoding system. A single character can be encoded in many different ways.

ASCII Characters and Encoding

- The American Standard Code for Information Interchange (ASCII) is a still widely used character encoding standard.
- Originally, ASCII encodes 128 specified characters into seven-bit natural numbers. Extended ASCII encodes the 128 specified characters into eight-bit natural numbers. This makes code points available for additional characters.
- ISO 8859 is a family of extended ASCII codes that support different language requirements, for example:
 - ISO 8859-1 adds characters for most common Western European languages
 - ISO 8859-2 adds characters for the most common Eastern European languages
 - ISO 8859-5 adds characters for Cyrillic languages
- Unfortunately, ISO 8859 code points overlap, making it difficult to represent texts requiring many different character sets.

ASCII Characters and Code Points (decimal)

0 nul	1 soh	2 stx	3 etx	4 eot	5 enq	6 ack	7 bel
8 bs	9 ht	10 nl	11 vt	12 np	13 cr	14 so	15 si
16 dle	17 dc1	18 dc2	19 dc3	20 dc4	21 nak	22 syn	23 etb
24 can	25 em	26 sub	27 esc	28 fs	29 gs	30 rs	31 us
32 sp	33 !	34 "	35 #	36 \$	37 %	38 &	39 '
40 (41)	42 *	43 +	44 ,	45 -	46 .	47 /
48 0	49 1	50 2	51 3	52 4	53 5	54 6	55 7
56 8	57 9	58 :	59 ;	60 <	61 =	62 >	63 ?
64 @	65 A	66 B	67 C	68 D	69 E	70 F	71 G
72 H	73 I	74 J	75 K	76 L	77 M	78 N	79 O
80 P	81 Q	82 R	83 S	84 T	85 U	86 V	87 W
88 X	89 Y	90 Z	91 [92 \	93]	94 ^	95 _
96 `	97 a	98 b	99 c	100 d	101 e	102 f	103 g
104 h	105 i	106 j	107 k	108 l	109 m	110 n	111 o
112 p	113 q	114 r	115 s	116 t	117 u	118 v	119 w
120 x	121 y	122 z	123 {	124	125 }	126 ~	127 del

Universal Coded Character Set and Unicode

- The Universal Coded Character Set (UCS) is a standard set of characters defined and maintained by the International Organization of Standardization (ISO).
- As of Unicode 10.0 (June 2017), it contains 136 690 abstract characters, each identified by an unambiguous name and an integer number called its code point.
- The Unicode Consortium produces industry standards based on the UCS for the encoding, representation, and handling of text expressed in most of the world's writing systems.
- Unicode can be implemented using different character encodings.
- The UTF-32 encoding encodes character code points directly into 32-bit numbers (fixed length encoding). While simple, an ASCII text of size n would become a UTF-32 text of size $4n$.

Unicode Transformation Format UTF-8

bytes	cp bits	first cp	last cp	byte 1	byte 2	bytes 3	byte 4
1	7	U+0000	U+007F	0xxxxxxx			
2	11	U+0080	U+07FF	110xxxxx	10xxxxxx		
3	16	U+0800	U+FFFF	1110xxxx	10xxxxxx	10xxxxxx	
4	21	U+10000	U+10FFFF	11110xxx	10xxxxxx	10xxxxxx	10xxxxxx

- Variable-length encoding of Unicode code points (cp) in such a way that seven-bit ASCII becomes valid UTF-8.
- The € symbol with the code point U+20AC (0010 0000 1010 1100 in binary notation) encodes as 0xE282AC (11100010 10000010 10101100 in binary notation).
- Note that this makes the € more expensive than the \$. ☺

- Let Σ be a non-empty finite set of symbols (or characters), called the alphabet.
- A string (or word) over Σ is any finite sequence of symbols from Σ , including (of course) the empty sequence.
- Typical operations on strings are `length()`, `concatenation()`, `reverse()`, ...
- There are different ways to store strings internally. Two common approaches are:
 - The sequence is *null-terminated*, i.e., the characters of the string are followed by a special NUL character.
 - The sequence is *length-prefixed*, i.e., a natural number indicating the length of the string is stored in front of the characters.
- In some programming languages, you need to know how strings are stored, in other languages you happily leave the details to the language implementation.

Date and Time

12 Natural Numbers

13 Integer Numbers

14 Rational and Real Numbers

15 Floating Point Numbers

16 International System of Units

17 Characters and Strings

18 Date and Time

System Time and Clocks

- Computer systems usually maintain a notion of *system time*. The term system time indicates that two different systems usually have a different notion of system time.
- System time is measured by a *system clock*, which is typically implemented as a simple count of the number of ticks that have transpired since some arbitrary starting date, called the epoch.
- Since internal counting mechanisms are not very precise, systems often exchange time information with other systems that have “better” clocks or sources of time in order to converge their notions of time.
- Time is sometimes used to order events, due to its monotonic nature.
- In distributed systems, this has its limitations and therefore the notion of logical clocks has been invented. (Logical clocks do not measure time, they only help to order events.)

Calendar Time

- System time can be converted into *calendar time*, a reference to a particular time represented within a calendar system.
- A popular calendar is the *Gregorian calendar*, which maps a time reference into a year, a month within the year, and a day within a month.
- The Gregorian calendar was introduced by Pope Gregory XIII in October 1582.
- The Coordinated Universal Time (UTC) is the primary time standard by which the world regulates clocks and time.
- Due to the rotation of the earth, days start and end at different moments. This is reflected by the notion of a *time zone*, which is essentially an offset to UTC.
- The number of time zones is not static and time zones change occasionally.

ISO 8601 Date and Time Formats

- Different parts of the world use different formats to write down a calendar time, which can easily lead to confusion.
- The ISO 8601 standard defines an unambiguous notation for calendar time.
- ISO 8601 in addition defines formats for durations and time intervals.

name	format	example
date	yyyy-mm-dd	2017-06-13
time	hh:mm:ss	15:22:36
date and time	yyyy-mm-ddThh:mm:ss[±hh:mm]	2017-06-13T15:22:36+02:00
date and time	yyyy-mm-ddThh:mm:ss[±hh:mm]	2017-06-13T13:22:36+00:00
date and time	yyyy-mm-ddThh:mm:ssZ	2017-06-13T13:22:36Z
date and week	yyyy-Www	2017-W24

Part: Boolean Algebra and Logic

- 19 Elementary Boolean Functions
- 20 Boolean Functions and Formulas
- 21 Boolean Algebra Equivalence Laws
- 22 Normal Forms (CNF and DNF)
- 23 Complexity of Boolean Formulas
- 24 Boolean Logic and the Satisfiability Problem

Logic can be confusing...

- If all men are mortal and Socrates is a man, then Socrates is mortal.
- I like Pat or I like Joe.
If I like Pat, I like Joe.
Do I like Joe?
- If cats are dogs, then the sun shines.
- “Logic is the beginning of wisdom, not the end of it.”

Elementary Boolean Functions

- 19 Elementary Boolean Functions
- 20 Boolean Functions and Formulas
- 21 Boolean Algebra Equivalence Laws
- 22 Normal Forms (CNF and DNF)
- 23 Complexity of Boolean Formulas
- 24 Boolean Logic and the Satisfiability Problem

Boolean Variables

- Boolean logic describes objects that can take only one of two values.
- The values may be different voltage levels $\{0, V^+\}$ or special symbols $\{F, T\}$ or simply the digits $\{0, 1\}$.
- In artificial intelligence, such objects are often called *propositions* and they are either *true* or *false*.
- In mathematics, the objects are called *Boolean variables* and we use the symbols X_1, X_2, X_3, \dots for them.
- The main purpose of Boolean logic is to describe (or design) interdependencies between Boolean variables.

Interpretation of Boolean Variables

Definition (Boolean variables)

A Boolean variable X_i with $i \geq 1$ is an object that can take on one of the two values 0 or 1. The set of all Boolean variables is $\mathbf{X} = \{X_1, X_2, X_3, \dots\}$.

Definition (Interpretation)

Let \mathbf{D} be a subset of \mathbf{X} . An *interpretation* \mathcal{I} of \mathbf{D} is a function $\mathcal{I} : \mathbf{D} \mapsto \{0, 1\}$.

- The set \mathbf{X} is very large. It is often sufficient to work with a suitable subset \mathbf{D} of \mathbf{X} .
- An interpretation assigns to every Boolean variable a value.
- An interpretation is also called a truth value assignment.

Boolean \wedge Function (and)

X	Y	$X \wedge Y$
0	0	0
0	1	0
1	0	0
1	1	1

- The logical *and* \wedge can be viewed as a function that maps two Boolean values to a Boolean value:

$$\wedge : \{0, 1\} \times \{0, 1\} \mapsto \{0, 1\}$$

- A truth table defines a Boolean operation (or function) by listing the result for all possible arguments.
- In programming languages like C or C++ (or even Haskell), the operator `&&` is often used to represent the \wedge operation.

Boolean \vee Function (or)

X	Y	$X \vee Y$
0	0	0
0	1	1
1	0	1
1	1	1

- The logical *or* \vee can be viewed as a function that maps two Boolean values to a Boolean value:

$$\vee : \{0, 1\} \times \{0, 1\} \mapsto \{0, 1\}$$

- Each row in the truth table corresponds to one interpretation.
 - A truth table simply lists all possible interpretations.
- In programming languages like C or C++ (or even Haskell), the operator `||` is often used to represent the \vee operation.

Boolean \neg Function (not)

X	$\neg X$
0	1
1	0

- The logical *not* \neg can be viewed as a unary function that maps a Boolean value to a Boolean value:

$$\neg : \{0, 1\} \mapsto \{0, 1\}$$

- The \neg operation simply flips a Boolean value.
- In programming languages like C or C++, the operator `!` is often used to represent the \neg operation.

Boolean \rightarrow Function (implies)

X	Y	$X \rightarrow Y$
0	0	1
0	1	1
1	0	0
1	1	1

- The logical *implication* \rightarrow can be viewed as a function that maps two Boolean values to a Boolean value:

$$\rightarrow: \{0, 1\} \times \{0, 1\} \mapsto \{0, 1\}$$

- The implication represents statements of the form “if X then Y ” (where X is called the precondition and Y the consequence).
- The logical implication is often confusing to ordinary mortals. A logical implication is false only if the precondition is true, but the consequence it asserts is false.
- The claim “if cats eat dogs, then the sun shines” is logically true.

Boolean \leftrightarrow Function (equivalence)

X	Y	$X \leftrightarrow Y$
0	0	1
0	1	0
1	0	0
1	1	1

- The logical *equivalence* \leftrightarrow can be viewed as a function that maps two Boolean values to a Boolean value:

$$\leftrightarrow: \{0, 1\} \times \{0, 1\} \mapsto \{0, 1\}$$

- In programming languages like C or C++, the operator `==` is often used to represent the equivalence function as an operation.

Boolean $\dot{\vee}$ Function (exclusive or)

X	Y	$X \dot{\vee} Y$
0	0	0
0	1	1
1	0	1
1	1	0

- The logical *exclusive or* $\dot{\vee}$ can be viewed as a function that maps two Boolean values to a Boolean value:

$$\dot{\vee} : \{0, 1\} \times \{0, 1\} \mapsto \{0, 1\}$$

- Another commonly used symbol for the exclusive or is \oplus .

Boolean \uparrow Function (not-and)

X	Y	$X \uparrow Y$
0	0	1
0	1	1
1	0	1
1	1	0

- The logical *not-and* (nand) or \uparrow can be viewed as a function that maps two Boolean values to a Boolean value:

$$\uparrow: \{0, 1\} \times \{0, 1\} \mapsto \{0, 1\}$$

- The \uparrow function is also called Sheffer stroke.
- While we use the functions \wedge , \vee , and \neg to introduce more complex Boolean functions, the Sheffer stroke is sufficient to derive all elementary Boolean functions from it.
- This is important for digital circuits since all you need are not-and gates.

Boolean \downarrow Function (not-or)

X	Y	$X \downarrow Y$
0	0	1
0	1	0
1	0	0
1	1	0

- The logical *not-or* (nor) \downarrow can be viewed as a function that maps two Boolean values to a Boolean value:

$$\downarrow: \{0, 1\} \times \{0, 1\} \mapsto \{0, 1\}$$

- The \downarrow function is also called Quine arrow.

- The \downarrow (nor) is like \uparrow (nand) sufficient to derive all elementary Boolean functions.

Boolean Functions and Formulas

- 19 Elementary Boolean Functions
- 20 Boolean Functions and Formulas**
- 21 Boolean Algebra Equivalence Laws
- 22 Normal Forms (CNF and DNF)
- 23 Complexity of Boolean Formulas
- 24 Boolean Logic and the Satisfiability Problem

Boolean Functions

- Elementary Boolean functions (\neg, \wedge, \vee) can be composed to define more complex functions.
- An example of a composed function is

$$\varphi(X, Y) := \neg(X \wedge Y)$$

which is a function $\varphi : \{0, 1\} \times \{0, 1\} \mapsto \{0, 1\}$ and means “first compute the \wedge of X and Y , then apply the \neg on the result you got from the \wedge ”.

- Boolean functions can take a large number of arguments. Here is a function taking three arguments:

$$\varphi(X, Y, Z) := (\neg(X \wedge Y) \vee (Z \wedge Y))$$

- The left hand side of the notation above defines the function name and its arguments, the right hand side defines the function itself by means of a formula.

Definition (Boolean function)

A *Boolean function* φ is any function of the type $\varphi : \{0, 1\}^k \mapsto \{0, 1\}$, where $k \geq 0$.

- The number k of arguments is called the arity of the function.
- A Boolean function with arity $k = 0$ assigns truth values to nothing. There are two such functions, one always returning 0 and the other always returning 1. We simply identify these two arity-0 functions with the truth value constants 0 and 1.
- The truth table of a Boolean function with arity k has 2^k rows. For a function with a large arity, truth tables become quickly unmanageable.

Syntax of Boolean formulas (aka Boolean expressions)

Definition (Syntax of Boolean formulas)

Basis of inductive definition:

- 1a Every Boolean variable X_i is a Boolean formula.
- 1b The two Boolean constants 0 and 1 are Boolean formulas.

Induction step:

- 2a If φ and ψ are Boolean formulas, then $(\varphi \wedge \psi)$ is a Boolean formula.
- 2b If φ and ψ are Boolean formulas, then $(\varphi \vee \psi)$ is a Boolean formula.
- 2c If φ is a Boolean formula, then $\neg\varphi$ is a Boolean formula.

Semantics of Boolean formulas

Definition (Semantics of Boolean formulas)

Let $\mathbf{D} \subseteq \mathbf{X}$ be a set of Boolean variables and $\mathcal{I} : \mathbf{D} \mapsto \{0, 1\}$ an interpretation. Let $\Phi(\mathbf{D})$ be the set of all Boolean formulas which contain only Boolean variables that are in \mathbf{D} . We define a generalized version of an interpretation $\mathcal{I}^* : \Phi(\mathbf{D}) \mapsto \{0, 1\}$.

Basis of the inductive definition:

- 1a For every Boolean variable $X \in \mathbf{D}$, $\mathcal{I}^*(X) = \mathcal{I}(X)$.
- 1b For the two Boolean constants 0 and 1, we set $\mathcal{I}^*(0) = 0$ and $\mathcal{I}^*(1) = 1$.

Definition (Semantics of Boolean formulas (cont.))

Induction step, with φ and ψ in $\Phi(\mathbf{D})$:

2a

$$\mathcal{I}^*((\varphi \wedge \psi)) = \begin{cases} 1 & \text{if } \mathcal{I}^*(\varphi) = 1 \text{ and } \mathcal{I}^*(\psi) = 1 \\ 0 & \text{otherwise} \end{cases}$$

2b

$$\mathcal{I}^*((\varphi \vee \psi)) = \begin{cases} 1 & \text{if } \mathcal{I}^*(\varphi) = 1 \text{ or } \mathcal{I}^*(\psi) = 1 \\ 0 & \text{otherwise} \end{cases}$$

2c

$$\mathcal{I}^*(\neg\varphi) = \begin{cases} 1 & \text{if } \mathcal{I}^*(\varphi) = 0 \\ 0 & \text{if } \mathcal{I}^*(\varphi) = 1 \end{cases}$$

Boolean Algebra Equivalence Laws

- 19 Elementary Boolean Functions
- 20 Boolean Functions and Formulas
- 21 Boolean Algebra Equivalence Laws**
- 22 Normal Forms (CNF and DNF)
- 23 Complexity of Boolean Formulas
- 24 Boolean Logic and the Satisfiability Problem

Tautology and contradiction

Definition (adapted interpretation)

An interpretation $\mathcal{I} : \mathbf{D} \mapsto \{0, 1\}$ is *adapted* to a Boolean formula φ if all Boolean variables that occur in φ are contained in \mathbf{D} .

Definition (tautologies and contradictions)

A Boolean formula φ is a *tautology* if for all interpretations \mathcal{I} which are adapted to φ it holds that $\mathcal{I}(\varphi) = 1$. A Boolean formula is a *contradiction* if for all interpretations \mathcal{I} which are adapted to φ it holds that $\mathcal{I}(\varphi) = 0$.

Satisfying a Boolean formula

Definition (satisfying a Boolean formula)

An interpretation \mathcal{I} which is adapted to a Boolean formula φ is said to *satisfy* the formula φ if $\mathcal{I}(\varphi) = 1$. A formula φ is called *satisfiable* if there exists an interpretation which satisfies φ .

The following two statements are equivalent characterizations of satisfiability:

- A Boolean formula is satisfiable if and only if its truth table contains at least one row that results in 1.
- A Boolean formula is satisfiable if and only if it is not a contradiction.

Equivalence of Boolean formulas

Definition (equivalence of Boolean formulas)

Let φ , ψ be two Boolean formulas. The formula φ is equivalent to the formula ψ , written $\varphi \equiv \psi$, if for all interpretations \mathcal{I} which are adapted to both φ and ψ it holds that $\mathcal{I}(\varphi) = \mathcal{I}(\psi)$.

- There are numerous “laws” of Boolean logic which are stated as equivalences. Each of these laws can be proven by writing down the corresponding truth table.
- Boolean equivalence “laws” can be used to “calculate” with logics, executing stepwise transformations from a starting formula to some target formula, where each step applies one equivalence law.

Equivalence laws

Proposition (equivalence laws)

For any Boolean formulas φ, ψ, χ , the following equivalences hold:

1. $\varphi \wedge 1 \equiv \varphi$ *and* $\varphi \vee 0 \equiv \varphi$ (identity)
2. $\varphi \vee 1 \equiv 1$ *and* $\varphi \wedge 0 \equiv 0$ (domination)
3. $(\varphi \wedge \varphi) \equiv \varphi$ *and* $(\varphi \vee \varphi) \equiv \varphi$ (idempotency)
4. $(\varphi \wedge \psi) \equiv (\psi \wedge \varphi)$ *and* $(\varphi \vee \psi) \equiv (\psi \vee \varphi)$ (commutativity)
5. $((\varphi \wedge \psi) \wedge \chi) \equiv (\varphi \wedge (\psi \wedge \chi))$ *and* $((\varphi \vee \psi) \vee \chi) \equiv (\varphi \vee (\psi \vee \chi))$ (associativity)
6. $\varphi \wedge (\psi \vee \chi) \equiv (\varphi \wedge \psi) \vee (\varphi \wedge \chi)$ *and* $\varphi \vee (\psi \wedge \chi) \equiv (\varphi \vee \psi) \wedge (\varphi \vee \chi)$ (distributivity)
7. $\neg\neg\varphi \equiv \varphi$ (double negation)
8. $\neg(\varphi \wedge \psi) \equiv (\neg\varphi \vee \neg\psi)$ *and* $\neg(\varphi \vee \psi) \equiv (\neg\varphi \wedge \neg\psi)$ (de Morgan's laws)
9. $\varphi \wedge (\varphi \vee \psi) \equiv \varphi$ *and* $\varphi \vee (\varphi \wedge \psi) \equiv \varphi$ (absorption laws)

Normal Forms (CNF and DNF)

- 19 Elementary Boolean Functions
- 20 Boolean Functions and Formulas
- 21 Boolean Algebra Equivalence Laws
- 22 Normal Forms (CNF and DNF)**
- 23 Complexity of Boolean Formulas
- 24 Boolean Logic and the Satisfiability Problem

Literals

Definition (literals)

A *literal* L_i is a Boolean formula that has one of the forms X_i , $\neg X_i$, 0 , 1 , $\neg 0$, $\neg 1$, i.e., a literal is either a Boolean variable or a constant or a negation of a Boolean variable or a constant. The literals X_i , 0 , 1 are called *positive literals* and the literals $\neg X_i$, $\neg 0$, $\neg 1$ are called *negative literals*.

Definition (monomial)

A *monomial* (or *product term*) is a literal or the logic and (product) of literals.

Definition (clause)

A *clause* (or *sum term*) is a literal or the logic or (sum) of literals.

Conjunctive Normal Form

Definition (conjunctive normal form)

A Boolean formula is said to be in *conjunctive normal form* (CNF) if it is a conjunction of disjunctions of literals.

- Examples of formulas in CNF:

- X_1

(this is a short form of $(1 \vee 1) \wedge (X_1 \vee 0)$)

- $X_1 \wedge X_2$

(this is a short form of $(X_1 \vee X_1) \wedge (X_2 \vee X_2)$)

- $X_1 \vee X_2$

(this is a short form of $(1 \vee 1) \wedge (X_1 \vee X_2)$)

- $\neg X_1 \wedge (X_2 \vee X_3)$

(this is a short form of $(0 \vee \neg X_1) \wedge (X_2 \vee X_3)$)

- $(X_1 \vee \neg X_2) \wedge (\neg X_1 \vee X_2)$

- We typically write the short form, leaving out trivial expansions into full CNF form.

Disjunctive Normal Form

Definition (disjunctive normal form)

A Boolean formula is said to be in *disjunctive normal form* (DNF) if it is a disjunction of conjunctions of literals.

- Examples of formulas in DNF:

- X_1 (this is a short form of $(0 \wedge 0) \vee (X_1 \wedge 1)$)
- $X_1 \wedge X_2$ (this is a short form of $(0 \wedge 0) \vee (X_1 \wedge X_2)$)
- $X_1 \vee X_2$ (this is a short form of $(X_1 \wedge X_1) \vee (X_2 \wedge X_2)$)
- $(\neg X_1 \wedge X_2) \vee (\neg X_1 \wedge X_3)$
- $(\neg X_1 \wedge \neg X_2) \vee (X_1 \wedge X_2)$

- We typically write the short form, leaving out trivial expansions into full DNF form.

Equivalence of Normal Forms

Proposition (CNF equivalence)

Every Boolean formula φ is equivalent to a Boolean formula χ in conjunctive normal form.

Proposition (DNF equivalence)

Every Boolean formula φ is equivalent to a Boolean formula χ in disjunctive normal form.

- These two results are important since we can represent any Boolean formula in a “shallow” format that does not need any “deeply nested” bracketing levels.

Obtaining a DNF from a Truth Table

- Given a truth table, a DNF can be obtained by writing down a conjunction of the input values for every row where the result is 1 and connecting all obtained conjunctions together with a disjunction.

X	Y	$X \vee Y$
0	0	0
0	1	1
1	0	1
1	1	0

- 2nd row: $\neg X \wedge Y$
- 3rd row: $X \wedge \neg Y$
- $\chi = (\neg X \wedge Y) \vee (X \wedge \neg Y)$

Obtaining a CNF from a Truth Table

- Given a truth table, a CNF can be obtained by writing down a disjunction of the negated input values for every row where the result is 0 and connecting all obtained disjunctions together with a conjunction.

X	Y	$X \vee Y$
0	0	0
0	1	1
1	0	1
1	1	0

- 1st row: $X \vee Y$
- 4th row: $\neg X \vee \neg Y$
- $\chi = (X \vee Y) \wedge (\neg X \vee \neg Y)$

Complexity of Boolean Formulas

- 19 Elementary Boolean Functions
- 20 Boolean Functions and Formulas
- 21 Boolean Algebra Equivalence Laws
- 22 Normal Forms (CNF and DNF)
- 23 Complexity of Boolean Formulas**
- 24 Boolean Logic and the Satisfiability Problem

Cost of Boolean Expressions and Functions

Definition (cost of boolean expression)

The cost $C(e)$ of a boolean expression e is the number of operators in e .

Definition (cost of boolean function)

The cost $C(f)$ of a boolean function f is the minimum cost of boolean expressions defining f , $C(f) = \min\{C(e) | e \text{ defines } f\}$.

- We can find expressions of arbitrary high cost for a given boolean function.
- How do we find an expression with minimal cost for a given boolean function?

Implicants and Prime Implicants

Definition (implicant)

A product term P of a Boolean function φ of n variables is called an *implicant* of φ if and only if for every combination of values of the n variables for which P is true, φ is also true.

Definition (prime implicant)

An implicant of a function φ is called a *prime implicant* of φ if it is no longer an implicant if any literal is deleted from it.

Definition (essential prime implicant)

A prime implicant of a function φ is called an *essential prime implicant* of φ if it covers a true output of φ that no combination of other prime implicants is covers.

Quine McCluskey Algorithm

- QM-0 Find all implicants of a given function (e.g., by determining the DNF from a truth table or by converting a boolean expression into DNF).
- QM-1 Repeatedly combine non-prime implicants until there are only prime implicants left.
- QM-2 Determine a minimum sum of prime implicants that defines the function. (This sum not necessarily includes all prime implicants.)
- We will further detail the steps QM-1 and QM-2 in the following slides.
 - See also the complete example in the notes.

Finding Prime Implicants (QM-1)

- PI-1 Classify and sort the minterms by the number of positive literals they contain.
- PI-2 Iterate over the classes and compare each minterms of a class with all minterms of the following class. For each pair that differs only in one bit position, mark the bit position as a wildcard and write down the newly created shorter term combining two terms. Mark the two terms as used.
- PI-3 Repeat the last step if new combined terms were created.
- PI-4 The set of minterms or combined terms not marked as used are the prime implicants.

Finding Minimal Sets of Prime Implicants (QM-2)

MS-1 Identify essential prime implications (essential prime implicants cover an implicant that is not covered by any of the other prime implicants)

MS-2 Find a minimum coverage of the remaining implicants by the remaining prime implicants

- Note that multiple minimal coverages may exist. The algorithm above does not define which solution is returned in this case.
- There are additional ways to cut the search space by eliminating rows or columns that are “dominated” by other rows or columns.

Boolean Logic and the Satisfiability Problem

- 19 Elementary Boolean Functions
- 20 Boolean Functions and Formulas
- 21 Boolean Algebra Equivalence Laws
- 22 Normal Forms (CNF and DNF)
- 23 Complexity of Boolean Formulas
- 24 Boolean Logic and the Satisfiability Problem

Logic Statements

- A common task is to decide whether statements of the form
if *premises* P_1 **and** ... **and** P_m *hold*, **then** *conclusion* C *holds*
are true.
- The premises P_i and the conclusion C are expressed in some logic formalism, the simplest is Boolean logic (also called propositional logic).
- Restricting us to Boolean logic here, the statement above can be seen as a Boolean formula of the following structure

$$(\varphi_1 \wedge \dots \wedge \varphi_m) \rightarrow \psi$$

and we are interested to find out whether such a formula is true, i.e., whether it is a tautology.

Tautology and Satisfiability

- Recall that a Boolean formula τ is a tautology if and only if $\tau' = \neg\tau$ is a contradiction. Furthermore, a Boolean formula is a contradiction if and only if it is not satisfiable. Hence, in order to check whether

$$\tau = (\varphi_1 \wedge \dots \wedge \varphi_m) \rightarrow \psi \tag{1}$$

is a tautology, we may check whether

$$\tau' = \neg((\varphi_1 \wedge \dots \wedge \varphi_m) \rightarrow \psi) \tag{2}$$

is unsatisfiable.

- If we show that τ' is satisfiable, we have disproven τ .

Tautology and Satisfiability

- Since $\varphi \rightarrow \psi \equiv \neg(\varphi \wedge \neg\psi)$, we can rewrite the formulas as follows:

$$\tau = (\varphi_1 \wedge \dots \wedge \varphi_m) \rightarrow \psi = \neg(\varphi_1 \wedge \dots \wedge \varphi_m \wedge \neg\psi) \quad (3)$$

$$\tau' = \neg((\varphi_1 \wedge \dots \wedge \varphi_m) \rightarrow \psi) = (\varphi_1 \wedge \dots \wedge \varphi_m \wedge \neg\psi) \quad (4)$$

- To disprove τ , it is often easier to prove that τ' is satisfiable.
- Note that τ' has a homogenous structure. If we transform the elements $\varphi_1, \dots, \varphi_m, \psi$ into CNF, then the entire formula is in CNF.
- If τ' is in CNF, all we need is to invoke an algorithm that searches for interpretations \mathcal{I} which satisfy a formula in CNF. If there is such an interpretation, τ is disproven, otherwise, if there is no such interpretation, then τ is proven.

Satisfiability Problem

Definition (satisfiability problem)

The satisfiability problem (SAT) is the following computational problem: Given as input a Boolean formula in CNF, compute as output a “yes” or “no” response according to whether the input formula is satisfiable or not.

- It is believed that there is no polynomial time solution for this problem.

Part: Computer Architecture and System Software

25 Logic Gates and Digital Circuits

26 Von Neumann Computer Architecture

27 Interpreter and Compiler

28 Operating Systems

Logic Gates and Digital Circuits

25 Logic Gates and Digital Circuits

26 Von Neumann Computer Architecture

27 Interpreter and Compiler

28 Operating Systems

Recall elementary boolean operations and functions

- Recall the elementary boolean operations AND (\wedge), OR (\vee), and NOT (\neg) as well as the boolean functions XOR ($\dot{\vee}$), NAND (\uparrow), and NOR (\downarrow).

$$X \dot{\vee} Y := (X \vee Y) \wedge \neg(X \wedge Y)$$


$$X \uparrow Y := \neg(X \wedge Y)$$

$$X \downarrow Y := \neg(X \vee Y)$$

- For each of these elementary boolean operations or functions, we can construct digital gates, for example, using transistors in Transistor-Transistor Logic (TTL).
- Note: NAND and NOR gates are called *universal gates* since all other gates can be constructed by using multiple NAND or NOR gates.

Logic gates implementing logic functions

NOT (\neg)



AND (\wedge)



OR (\vee)



XOR ($\dot{\vee}$)



NAND (\uparrow)



NOR (\downarrow)



- There are different sets of symbols for logic gates (do not get confused if you look into other sources of information).
- The symbols used here are the ANSI (American National Standards Institute) symbols.

Addition of decimal and binary numbers

$$\begin{array}{r} 2 \quad 0010 \\ + 5 \quad + 0101 \\ \hline 7 \quad 0111 \end{array}$$

$$\begin{array}{r} 3 \quad 0011 \\ + 3 \quad + 0011 \\ \quad 11 \\ \hline 6 \quad 0110 \end{array}$$


$$\begin{array}{r} 8 \quad 1000 \\ + 3 \quad + 0011 \\ \quad 1 \\ \hline 11 \quad 1011 \end{array}$$

- We are used to add numbers in the decimal number system.
- Adding binary numbers is essentially the same, except that we only have the digits 0 and 1 at our disposal and “carry overs” are much more frequent.

Adding two bits (half adder)

- The half adder adds two single binary digits A and B .
- It has two outputs, sum (S) and carry (C).

A	B	C	S
0	0	0	0
0	1	0	1
1	0	0	1
1	1	1	0




Adding two bits (full adder)

- A full adder adds two single bit digits A and B and accounts for a carry bit C_{in} .
- It has two outputs, sum (S) and carry (C_{out}).

A	B	C_{in}	C_{out}	S
0	0	0	0	0
0	1	0	0	1
1	0	0	0	1
1	1	0	1	0
0	0	1	0	1
0	1	1	1	0
1	0	1	1	0
1	1	1	1	1


$$S = A \dot{\vee} B \dot{\vee} C_{in}$$

$$C_{out} = (A \wedge B) \vee (C_{in} \wedge (A \dot{\vee} B))$$



Adding N bits (ripple carry adder)

- And N-bit adder can be created using multiple full adders.
- Each full adder inputs a C_{in} , which is the C_{out} of the previous adder.
- Each carry bit “ripples” to the next full adder.



Von Neumann Computer Architecture


25 Logic Gates and Digital Circuits

26 Von Neumann Computer Architecture

27 Interpreter and Compiler


28 Operating Systems

Von Neumann computer architecture




- Control unit contains an instruction register and a program counter
- Arithmetic/logic unit (ALU) performs integer arithmetic and logical operations
- Program instructions and data is stored in a memory unit
- Processor registers provide small amount of storage as part of a central processing unit
- The central processing unit (CPU) carries out the actual computations

Computer system bus (data, address, and control)



- The *data bus* transports data (primarily between registers and main memory).
- The *address bus* selects which memory cell is being read or written.
- The *control bus* activates components and steers the data flow over the data bus and the usage of the address bus.

Simple Central Processing Unit



- Real CPUs usually have multiple registers
- Real CPUs support memory outside of the CPU itself
- Real CPUs have different instruction sets for different privilege levels
- Real CPUs have special digital circuits for floating point arithmetic or cryptographic operations

Instruction cycle (fetch – decode – execute cycle)

```
while True:  
    advance_program_counter();  
    instruction = fetch();  
    decode(instruction);  
    execute(instruction);
```

- The CPU runs in an endless loop fetching instructions, decoding them, and executing them.
- The set of instructions a CPU can execute is called the CPU's machine language
- Typical instructions are to add two N-bit numbers, to test whether a certain register is zero, or to jump to a certain position in the ordered list of machine instructions.
- An assembly programming language is a mnemonic representation of machine code.

Simple Machine Language

Op-code	Mnemonic	Function
001	LOAD	Load the value of the operand into the accumulator
010	STORE	Store the value of the accumulator at the address specified by the operand
011	ADD	Add the value of the operand to the accumulator
100	SUB	Subtract the value of the operand from the accumulator
101	EQUAL	If the value of the operand equals the value of the Accumulator, skip the next instruction
110	JUMP	Jump to a specified instruction by setting the program counter to the value of the operand
111	HALT	Stop execution

- Each instruction of the machine language is encoded into 8 bits:
 - 3 bits are used for the op-code
 - 1 bit is used to indicate whether the operand is a constant or a memory address
 - 4 bits are used to carry a constant or a memory address (the operand)

Program #1 in our simple machine language

#	Machine Code	Assembly Code	Description
0	001 1 0010	LOAD #2	Load the value 2 into the accumulator
1	010 0 1101	STORE 13	Store the value of the accumulator in memory location 13
2	001 1 0101	LOAD #5	Load the value 5 into the accumulator
3	010 0 1110	STORE 14	Store the value of the accumulator in memory location 14
4	001 0 1101	LOAD 13	Load the value of memory location 13 into the accumulator
5	011 0 1110	ADD 14	Add the value of memory location 14 to the accumulator
6	010 0 1111	STORE 15	Store the value of the accumulator in memory location 15
7	111 0 0000	HALT	Stop execution

- An animation of the execution of this program can be found here:
<http://courses.cs.vt.edu/csonline/MachineArchitecture/Lessons/CPU/Lesson.html>
- What is the equivalent C program?

Program #2 in our simple machine language

#	Machine Code	Assembly Code	Description
0	001 1 0101	LOAD #5	Load the value 5 into the accumulator
1	010 0 1111	STORE 15	Store the value of the accumulator in memory location 15
2	001 1 0000	LOAD #0	Load the value 0 into the accumulator
3	101 0 1111	EQUAL 15	Skip next instruction if accumulator equal to memory location 15
4	110 1 0110	JUMP #6	Jump to instruction 6 (set program counter to 6)
5	111 0 0000	HALT	Stop execution
6	011 1 0001	ADD #1	Add the value 1 to the accumulator
7	110 1 0011	JUMP #3	Jump to instruction 3 (set program counter to 3)

- An animation of the execution of this program can be found here:
<http://courses.cs.vt.edu/csonline/MachineArchitecture/Lessons/CPU/Lesson.html>
- What is the equivalent C program?

Interpreter and Compiler

25 Logic Gates and Digital Circuits

26 Von Neumann Computer Architecture

27 Interpreter and Compiler

28 Operating Systems

Are there better ways to write machine or assembler code?

- Observations:
 - Writing machine code or assembler code is difficult and time consuming.
 - Maintaining machine code or assembler code is even more difficult and time consuming (and most cost is spent on software maintenance).
- A high-level programming language is a programming language with strong abstraction from the low-level details of the computer.
- Rather than dealing with registers and memory addresses, high-level languages deal with variables, arrays, objects, collections, complex arithmetic or boolean expressions, subroutines and functions, loops, threads, locks, and other abstract computer science concepts, with a focus on usability over optimal program efficiency.

Simple C program to add two numbers

```
/* C source code
```

```
(C is a compiled procedural programming language) */
```

```
int main()
```

```
{
```

```
    int a = 5;
```

```
    int b = 2;
```

```
    int c = a + b;
```

```
    return c;
```

```
}
```

Disassembled machine code (without optimizations)

```
# compile without optimization (gcc) and look at the machine code
# gcc (Debian 4.7.2-5) 4.7.2 on Linux
```

```
00000000004004ac <main>:
```

4004ac:	55	push	%rbp
4004ad:	48 89 e5	mov	%rsp,%rbp
4004b0:	c7 45 fc 05 00 00 00	movl	\$0x5,-0x4(%rbp)
4004b7:	c7 45 f8 02 00 00 00	movl	\$0x2,-0x8(%rbp)
4004be:	8b 45 f8	mov	-0x8(%rbp),%eax
4004c1:	8b 55 fc	mov	-0x4(%rbp),%edx
4004c4:	01 d0	add	%edx,%eax
4004c6:	89 45 f4	mov	%eax,-0xc(%rbp)
4004c9:	8b 45 f4	mov	-0xc(%rbp),%eax
4004cc:	5d	pop	%rbp
4004cd:	c3	retq	
4004ce:	90	nop	
4004cf:	90	nop	

Disassembled machine code (with optimizations)

```
# compile with optimization (gcc -O2) and look at the machine code  
# gcc (Debian 4.7.2-5) 4.7.2 on Linux
```

```
00000000004003a0 <main>:
```

4003a0:	b8 07 00 00 00	mov	\$0x7,%eax
4003a5:	c3	retq	
4003a6:	90	nop	
4003a7:	90	nop	

Compiler and Interpreter

[1] Source Code --> Interpreter

[2] Source Code --> Compiler --> Machine Code

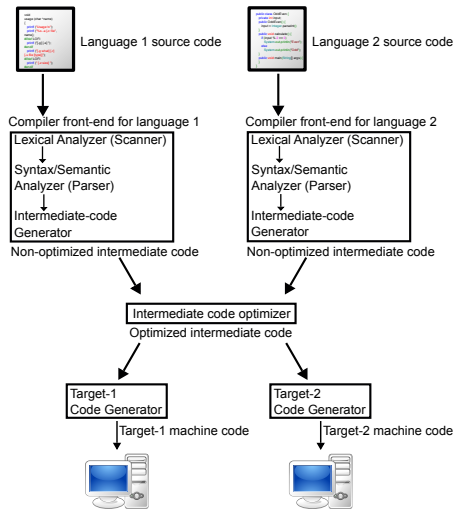
[3] Source Code --> Compiler --> Byte Code --> Interpreter

[4] Source Code --> Compiler --> Byte Code --> Compiler --> Machine Code

- An interpreter is a computer program that directly executes source code written in a higher-level programming language.
- A compiler is a program that transforms source code written in a higher-level programming language (the source language) into a lower-level computer language (the target language).

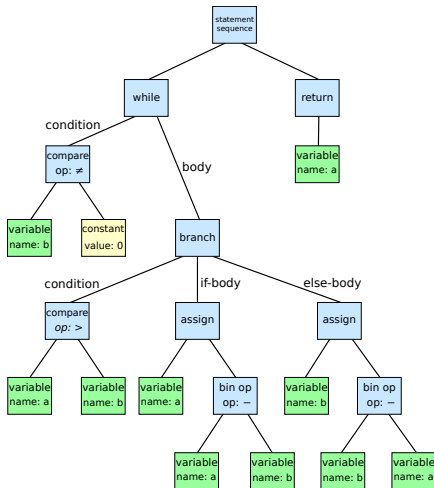
- A basic interpreter parses a statement, executes it, and moves on to the next statement (very similar to a fetch-decode-execute cycle).
- More advanced interpreter do a syntactic analysis to determine syntactic correctness before execution starts.
- Properties:
 - Highly interactive code development (trial-and-error coding)
 - Limited error detection capabilities before code execution starts
 - Interpretation causes a certain runtime overhead
 - Development of short pieces of code can be very fast
- Examples: command interpreter (shells), scripting languages

Compiler



- lexical analysis
⇒ sequence of token
- syntax analysis
⇒ parse tree
- semantic analysis
⇒ abstract syntax tree
- optimization
⇒ enhanced abstract syntax tree
- code generation
⇒ object code

Abstract Syntax Tree Example



Euclidean algorithm to find the greatest common divisor of a and b:

```
while (b != 0):
    if (a > b):
        a = a - b
    else:
        b = b - a
return a
```

Backus-Naur-Form and Formal Languages

The syntax of programming languages is often defined using syntax rules. A common notation for syntax rules is the Backus-Naur-Form (BNF):

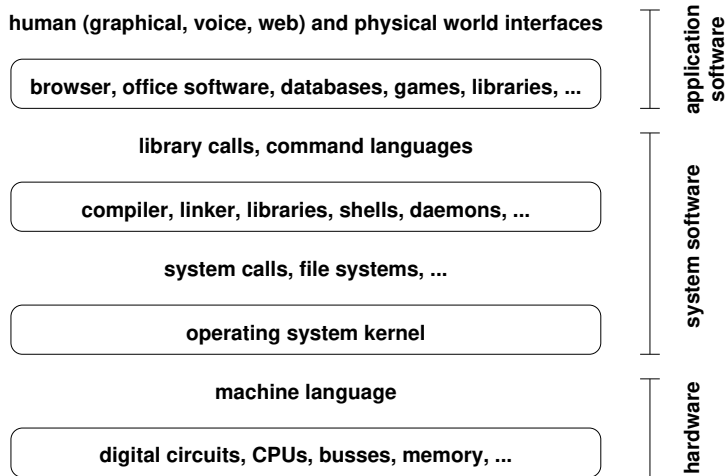
- Terminal symbols are enclosed in quotes
- Non-terminal symbols are enclosed in $\langle \rangle$
- A BNF rule consists of a non-terminal symbol followed by the defined-as operator $::=$ and a rule expression
- A rule expression consists of terminal and non-terminal symbols and operators; the empty operator denotes contatenation and the $|$ operator denotes an alternative
- Parenthesis may be used to group elements of a rule expression

A set of BNF rules has a non-terminal starting symbol.

Virtual Machines and Emulators

- A virtual machine (VM) is an emulation of a particular computer system. Virtual machines operate based on the computer architecture and functions of a real computer.
 - An emulator is hardware or software or both that duplicates (or emulates) the functions of one computer system (the guest) in another computer system (the host), different from the first one, so that the emulated behavior closely resembles the behavior of the real system (the guest).
- ⇒ Virtual machines were invented in the 1970s and reinvented in the 1990s.
- ⇒ Virtual machines have been an enabler for cloud computing since they are easy to start / stop / clone / migrate and they separate the software implementing services from the underlying hardware.

Hardware vs. System Software vs. Application Software



- 25 Logic Gates and Digital Circuits
- 26 Von Neumann Computer Architecture
- 27 Interpreter and Compiler
- 28 Operating Systems**

Operating System Kernel Functions

- Execute many programs concurrently (instead of just one program at a time)
- Assign resources to running programs (memory, CPU time, ...)
- Ensure a proper separation of concurrent processes
- Enforce resource limits and provide means to control processes
- Provide logical filesystems on top of block-oriented raw storage devices
- Control and coordinate input/output devices (keyboard, display, ...)
- Provide basic network communication services to applications
- Provide input/output abstractions that hide device specifics
- Enforce access control rules and privilege separation
- Provide a well defined application programming interface (API)

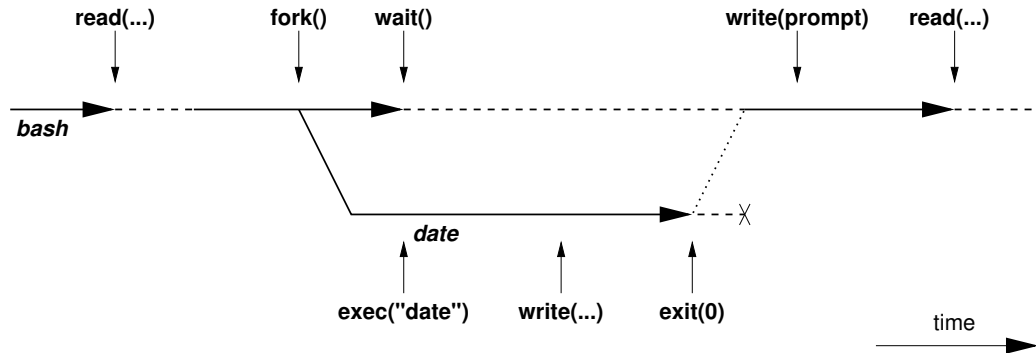
OS Abstraction #1: Processes and Process Lifecycle

Definition (process)

An instance of a computer program that is being executed is called a *process*.

- The OS kernel maintains information about each running process and assigns resources and ensures protection of concurrently running processes.
 - In Unix-like Operating Systems
 - a new process is created by “cloning” (forking) an already existing process
 - a process may load a new program image (machine code) to execute
 - a terminating process returns a number to its parent process
 - a parent process can wait for child processes to terminate
- ⇒ A very basic command interpreter can be written in a few lines of Python code.

OS Abstraction #1: Processes and Process Lifecycle




OS Abstraction #1: Processes and Process Lifecycle

```
while (1) {  
    show_prompt();           /* display prompt */  
    read_command();          /* read and parse command */  
    pid = fork();             /* create new process */  
    if (pid < 0) {            /* continue if fork() failed */  
        perror("fork");  
        continue;  
    }  
    if (pid != 0) {           /* parent process */  
        waitpid(pid, &status, 0); /* wait for child to terminate */  
    } else {                  /* child process */  
        execvp(args[0], args, 0); /* execute command */  
        perror("execvp");         /* only reach on exec failure */  
        _exit(1);                 /* exit without any cleanups */  
    }  
}
```

OS Abstraction #2: File Systems

- Files are persistent containers for the storage of data
- Unstructured files contain a sequence of bytes
- Applications interpret the content of a file in a specific way
- Files also have meta data (owner, permissions, timestamps)
- Hierarchical file systems use directories to organize files into a hierarchy
- Names of files and directories at one level of the hierarchy usually have to be unique
- The operating system maps the logical structure of a hierarchical file system to a block-oriented storage device
- The operating system must ensure file system integrity
- The operating system may support compression and encryption of file systems

OS Abstraction #2: File Systems (Unix)



- The logical structure of a typical Unix file system
- The `.` in a directory always refers to the directory itself
- The `..` in a directory always refers to the parent directory, except in the root directory
- A link is a reference of a file system object from a directory
- Any file system changes need to maintain the integrity of these links

OS Abstraction #2: File and Directory Operations (Unix)

File operations

<code>open()</code>	open a file
<code>read()</code>	read data from the current file position
<code>write()</code>	write data at the current file position
<code>seek()</code>	seek to a file position
<code>stat()</code>	read meta data
<code>close()</code>	close an open file
<code>unlink()</code>	remove a link to a file

Directory operations

<code>mkdir()</code>	create a directory
<code>rmdir()</code>	delete a directory
<code>chdir()</code>	change to a directory
<code>opendir()</code>	open a directory
<code>readdir()</code>	read a directory entry
<code>closedir()</code>	close a directory

OS Abstraction #3: Inter-process Communication

- Communication between processes:
 - Signals (software interrupts)
 - Pipes (local unidirectional byte streams)
 - Sockets (local and global bidirectional byte or datagram streams)
 - Shared memory (memory regions shared between multiple processes)
 - Message queues (a queue of messages between multiple processes)
 - ...
- Sockets are the basic inter-process communication abstraction used for communication between processes over the Internet

Part: Automata and Formal Languages

29 Finite State Machines

30 Pushdown Automaton

31 Turing Machines

32 Formal Languages

29 Finite State Machines

30 Pushdown Automaton

31 Turing Machines

32 Formal Languages

Definition (finite state machine)

A *finite state machine* (FSM) is a quintuple $(\Sigma, S, s_0, \delta, F)$ where:

- Σ is the input alphabet (a finite, non-empty set of symbols)
 - S is a finite, non-empty set of states
 - s_0 is an initial state ($s_0 \in S$)
 - δ is the state-transition function, $\delta : S \times \Sigma \mapsto S$
 - F is a possibly empty set of accepting states ($F \subset S$)
-
- We sometimes say that a FSM *accepts* an input word $w \in \Sigma^*$ if the machine, starting from the state s_0 , processes all symbols of w and reaches one of the accepting states in F .

Finite State Machine Example ($a^n b^m$)


- The FSM $(\Sigma, S, s_0, \delta, F)$ with $\Sigma = \{a, b\}$, $S = \{S0, S1, S2, S3\}$, $s_0 = S0$, $F = \{S2\}$, and

$$\delta = \{((S0, a), S1), ((S0, b), S3), \\ ((S1, a), S1), ((S1, b), S2), \\ ((S2, a), S3), ((S2, b), S2)\}$$

recognizes all words of the form $\{a^n b^m \mid n \geq 1, m \geq 1\}$.

- The set of words $w \in \Sigma^*$ accepted by a finite state is called the language recognized by the finite state machine.

FSM Example ($a^n b^m$) Represented as a Graph



- State machines can be represented as graphs where nodes (circles) represent states, arrows represent state transitions, an arrow pointing from a small black circle indicates the initial state, and accepting states are marked with a double circle.


Finite State Machine Example (integer)

- Consider the problem of deciding whether an input string contains an integer number, that is, whether it consists of at least one digit and an optional '-' at the very beginning (valid numbers would be -1, 0, -42, 42)
- The FSM $(\Sigma, S, s_0, \delta, F)$ with $\Sigma = \{-, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$, $S = \{S0, S1, S2, S3\}$, $s_0 = S0$, $F = \{S2\}$, and

$$\delta = \{((S0, -), S1), ((S0, 0), S2), \dots, ((S0, 9), S2), \\ ((S1, -), S3), ((S1, 0), S2), \dots, ((S1, 9), S2), \\ ((S2, -), S3), ((S2, 0), S2), \dots, ((S2, 9), S2), \\ ((S3, -), S3), ((S3, 0), S3), \dots, ((S3, 9), S3)\}$$

solves this integer number parsing problem.

FSM Example (integer) Represented as a Graph



Pushdown Automaton

29 Finite State Machines

30 Pushdown Automaton

31 Turing Machines

32 Formal Languages

Definition (pushdown automaton)

A *pushdown automaton* (PDA) is a 7-tuple $(\Sigma, S, s_0, \Gamma, Z, \delta, F)$ where:

- Σ is the input alphabet (a finite, non-empty set of symbols)
- S is a finite, non-empty set of states
- s_0 is an initial state ($s_0 \in S$)
- Γ is a finite, non-empty stack alphabet
- Z is the initial stack symbol ($Z \in \Gamma$)
- δ is the state-transition function, $\delta : S \times (\Sigma \cup \{\epsilon\}) \times \Gamma \mapsto S \times \Gamma^*$
- F is a possibly empty set of accepting states ($F \subset S$)


Pushdown Automaton Example ($a^n b^n$)

- The PDA $(\Sigma, S, s_0, \Gamma, Z, \delta)$ with $\Sigma = \{a, b\}$, $S = \{S0, S1, S2\}$, $s_0 = S0$, $\Gamma = \{A, Z\}$, $F = \{S2\}$, and

$$\delta = \{(S0, a, Z, S0, AZ), \\ (S0, a, A, S0, AA), \\ (S0, \epsilon, Z, S1, Z), \\ (S0, b, A, S1, \epsilon), \\ (S1, b, A, S1, \epsilon), \\ (S1, \epsilon, Z, S2, Z)\}$$

recognizes all words of the form $\{a^n b^n | n \geq 1\}$.

PDA Example ($a^n b^n$) Represented as a Graph



- Pushdown automata can be represented as graphs where nodes (circles) represent states, arrows represent state transitions, an arrow pointing from a small black circle indicates the initial state, and accepting states are marked with with a double circle.

Turing Machines

29 Finite State Machines

30 Pushdown Automaton

31 Turing Machines

32 Formal Languages

Definition (turing machine)

A *Turing machine* (TM) is a 7-tuple $(\Sigma, S, s_0, \Gamma, b, \delta, F)$ where:

- Σ is the set of input symbols ($\Sigma \subseteq \Gamma \setminus \{b\}$)
 - S is a finite, non-empty set of states
 - s_0 is an initial state ($s_0 \in S$)
 - Γ is a finite, non-empty set of tape alphabet symbols
 - b is the blank symbol ($b \in \Gamma$)
 - δ is the state-transition function, $\delta : (S \setminus F) \times \Gamma \mapsto S \times \Gamma \times \{L, R\}$
 - F is a set of accepting states ($F \subset S$)
- The symbol L indicates a left movement, the symbol R a right movement.


Turing Machine Example ($a^n b^n c^n$)

- The TM $(\Sigma, S, s_0, \Gamma, b, \delta, F)$ with $\Sigma = \{a, b\}$, $S = \{S0, S1, S2, S3, S4, S5\}$, $s_0 = S0$, $\Gamma = \{a, b, -\}$, $b = -$, $F = \{S5\}$, and

$$\begin{aligned} \delta = \{ & (S0, a, S1, A, R), (S0, B, S4, B, R), \\ & (S1, b, S2, B, R), (S1, a, S1, a, R), (S1, B, S1, B, R), \\ & (S2, c, S3, C, L), (S2, b, S2, b, R), (S2, C, S2, C, R), \\ & (S3, A, S0, A, R), (S3, C, S3, C, L), (S3, b, S3, b, L), \\ & (S3, B, S3, B, L), (S3, a, S3, a, L), \\ & (S4, -, S5, -, L), (S4, B, S4, B, R), (S4, C, S4, C, R) \} \end{aligned}$$

recognizes all words of the form $\{a^n b^n c^n | n \geq 1\}$.

TM Example ($a^n b^n c^n$) Represented as a Graph




Turing Machine Example (BB-3)

- The busy beaver game consists of designing a halting, binary-alphabet Turing machine, which writes the most 1s on the tape, using only a limited set of states.
- The TM $(\Sigma, S, s_0, \Gamma, b, \delta, F)$ with $\Sigma = \{1\}$, $S = \{S0, S1, S2, S3\}$, $s_0 = S0$, $\Gamma = \{0, 1\}$, $b = 0$, $F = \{S3\}$, and

$$\delta = \{(S0, 0, S1, 1, R), (S0, 1, S2, 1, L), \\ (S1, 0, S0, 1, L), (S1, 1, S1, 1, R), \\ (S2, 0, S1, 1, L), (S2, 1, S3, 1, R)\}$$

is a 3-state busy beaver (BB-3) Turing Machine (the halting state is not counted).

TM Example (BB-3) Represented as a Graph



- Turing machines can be represented as graphs where nodes (circles) represent states, arrows represent state transitions, an arrow pointing from a small black circle indicates the initial state, and accepting states are marked with a double circle.

Formal Languages

29 Finite State Machines

30 Pushdown Automaton

31 Turing Machines

32 Formal Languages

Formal Language and Formal Grammar

Definition (formal language)

Given an alphabet Σ , a *formal language* L is a subset of Σ^* , i.e., $\Sigma \subseteq L$. An element $w \in L$ is called a word of L .

Definition (formal grammar)

A *formal grammar* G is a tuple (N, Σ, P, S) where

- N is a finite set of non-terminal symbols (disjoint from Σ)
- Σ is a finite set of terminal symbols (disjoint from N)
- P is a finite set of production rules of the form $(\Sigma \cup N)^* N (\Sigma \cup N)^* \mapsto (\Sigma \cup N)^*$
- $S \in N$ is a distinguished start symbol

Definition (grammar derivation)

Given a formal grammar $G = (N, \Sigma, P, S)$, the binary relation \Rightarrow_G (pronounced as “ G derives in one step”) on strings in $(\Sigma \cup N)^*$ is defined by

$$x \Rightarrow_G y \text{ iff } \exists u, v, p, q \in (\Sigma \cup N)^* : (x = upv) \wedge (p \mapsto q \in P) \wedge (y = uqv).$$

Let \Rightarrow_G^* denote the reflexive transitive closure of \Rightarrow_G .

Definition (language of a grammar)

The language $L(G)$ of $G = (N, \Sigma, P, S)$ is defined as $\{w \in \Sigma^* \mid S \xRightarrow_G^* w\}$

Operations on Formal Languages

Definition (operations on languages)

Let L_1 and L_2 be formal languages. We defined the following operations:

- $L_1 \cup L_2 = \{w \mid w \in L_1 \vee w \in L_2\}$ (union)
- $L_1 \cap L_2 = \{w \mid w \in L_1 \wedge w \in L_2\}$ (intersection)
- $\bar{L}_1 = \{w \mid w \notin L_1\}$ (complement)
- $L_1 L_2 = \{wz \mid w \in L_1 \wedge z \in L_2\}$ (concatenation)
- $L_1^* = \{\epsilon\} \cup \{wz \mid w \in L_1 \wedge z \in L_1^*\}$ (Kleene star)

Definition (regular languages)

The collection of regular languages over an alphabet Σ is defined inductively as follows:

- The empty language \emptyset , and the empty string language $\{\epsilon\}$ are regular languages.
- For each $a \in \Sigma$, the singleton language $\{a\}$ is a regular language.
- If A and B are regular languages, then
 - $A \cup B$ (union),
 - AB (concatenation), and
 - A^* (Kleene star)are regular languages.
- No other languages over Σ are regular.

Regular Grammars

Definition (right regular grammar)

A formal grammar (N, Σ, P, S) is called a *right regular grammar* iff all the production rules in P are of one of the forms $A \mapsto a$, $A \mapsto aB$, or $A \mapsto \epsilon$ with $A, B \in N$, $a \in \Sigma$, and ϵ denoting the empty word.

Definition (left regular grammar)

A formal grammar (N, Σ, P, S) is called a *left regular grammar* iff all the production rules in P are of one of the forms $A \mapsto a$, $A \mapsto Ba$, or $A \mapsto \epsilon$ with $A, B \in N$, $a \in \Sigma$, and ϵ denoting the empty word.

Definition (regular grammar)

A *regular grammar* is a left or right regular grammar.

Properties of Regular Languages

- A regular language can be defined by a regular expression.
- A regular language can be accepted by a finite state machine.
- A regular language can be generated by a regular grammar.
- ...

⇒ For more properties of regular languages and proofs of the properties, see the course “Formal Languages and Logic”.

Context-Free Languages

Definition (context-free grammar)

A formal grammar (N, Σ, P, S) is called a *context-free grammar* iff all productions P are of the form $N \mapsto (\Sigma \cup N)^*$.

Definition (context-free language)

A *context-free language* is a formal language generated by a context-free grammar.

Properties of Context-Free Languages

- A context-free language can be accepted by a pushdown automata.
 - A context-free language can be generated by a context-free grammar.
 - The set of context-free languages includes the set of regular languages.
 - There are context-free languages that are not regular languages.
 - ...
- ⇒ For more properties of context-free languages and proofs of the properties, see the course “Formal Languages and Logic”.

Context-Sensitive Languages

Definition (context-sensitive grammar)

A formal grammar (N, Σ, P, S) is called a *context-sensitive grammar* iff all productions P are of the form $\alpha N \beta \mapsto \alpha \gamma \beta$ with $\alpha, \beta \in (\Sigma \cup N)^*$ and $\gamma \in (\Sigma \cup N)^+$

Definition (context-sensitive language)

A *context-sensitive language* is a formal language generated by a context-sensitive grammar.

Properties of Context-Sensitive Languages

- A context-sensitive language can be accepted by a Turing machine.
 - A context-sensitive language can be generated by a context-sensitive grammar.
 - The set of context-sensitive languages includes the set of context-free languages.
 - There are context-sensitive languages that are not context-free languages.
 - ...
- ⇒ For more properties of context-sensitive languages and proofs of the properties, see the course “Formal Languages and Logic”.

Part: Computability and Computational Complexity

33 Landau Sets and Big O Notation

34 Computability

35 Computational Complexity

Landau Sets and Big O Notation

33 Landau Sets and Big O Notation

34 Computability

35 Computational Complexity

Big O Notation (Landau Notation)

Definition (asymptotically bounded)

Let $f, g : \mathbb{N} \mapsto \mathbb{N}$ be two functions. We say that f is *asymptotically bounded* by g , written as $f \leq_a g$, if and only if there is an $n_0 \in \mathbb{N}$, such that $f(n) \leq g(n)$ for all $n > n_0$.

Definition (Landau Sets)

The three *Landau Sets* $O(g)$, $\Omega(g)$, $\Theta(g)$ are defined as follows:

- $O(g) = \{f | \exists k. f \leq_a k \cdot g\}$
- $\Omega(g) = \{f | \exists k. k \cdot g \leq_a f\}$
- $\Theta(g) = O(g) \cap \Omega(g)$

Commonly Used Landau Sets

Landau Set	class name	rank
$O(1)$	constant	1
$O(\log_2(n))$	logarithmic	2
$O(n)$	linear	3
$O(n \log_2(n))$	linear logarithmic	4

Landau Set	class name	rank
$O(n^2)$	quadratic	5
$O(n^k)$	polynomial	6
$O(k^n)$	exponential	7

Theorem (Landau Set Ranking)

The commonly used Landau Sets establish a ranking such that

$$O(1) \subset O(\log_2(n)) \subset O(n) \subset O(n \log_2(n)) \subset O(n^2) \subset O(n^k) \subset O(l^n)$$

for $k > 2$ and $l > 1$.

Theorem (Landau Set Computation Rules)

We have the following computation rules for Landau sets:

- *If $k \neq 0$ and $f \in O(g)$, then $(kf) \in O(g)$.*
- *If $f_1 \in O(g_1)$ and $f_2 \in O(g_2)$, then $(f_1 + f_2) \in O(|g_1| + |g_2|)$.*
- *If $f_1 \in O(g_1)$ and $f_2 \in O(g_2)$, then $(f_1 f_2) \in O(g_1 g_2)$.*

Examples:

- $f(n) = 42 \implies f \in O(1)$
- $f(n) = 26n + 72 \implies f \in O(n)$
- $f(n) = 856n^{10} + 123n^3 + 75 \implies f \in O(n^{10})$
- $f(n) = 3 \cdot 2^n + 42 \implies f \in O(2^n)$

Big O Notation (Usage)

- The Big O Notation describes the limiting behavior of a function when the argument tends towards a particular value of infinity.
- We classify a function describing the (time or space) complexity of an algorithm by determining the closest Landau Set it belongs to.
- Use O classes for worst case complexity, use Ω classes for best case complexity.

33 Landau Sets and Big O Notation

34 Computability

35 Computational Complexity

Turing Completeness and Equivalence

Definition (Turing complete)

A system of data-manipulation rules (such as a computer's instruction set, a programming language, or a cellular automaton) is said to be *Turing complete* or computationally universal if it can be used to simulate any Turing machine.

Definition (Turing equivalent)

Two computers P and Q are called *Turing equivalent* if P can simulate Q and Q can simulate P .

Definition (Church-Turing thesis)

The *Church-Turing thesis* states that a function on the natural numbers is computable by a human being following an algorithm, ignoring resource limitations, if and only if it is computable by a Turing machine.

- The Church-Turing thesis is a conjecture.

Halting Problem

Definition (halting problem)

The halting problem is a decision problem: Given an arbitrary program and an input to the program, whether the program will eventually halt when run with that input.

- It is impossible to decide the halting problem. The proof uses a diagonalization argument. Here is an outline of the basic idea...
- Assume that the halting problem for any program can be solved by a machine H .
- Using H , construct a machine G that goes into an endless loop if H determines that an algorithm halts.
- Feed the program G as input to G :
 - If G halts, then H decided that G does not halt, which is a contradiction.
 - If G does not halt, then H decided the G does halt, which is a contradiction.

Computational Complexity

33 Landau Sets and Big O Notation

34 Computability

35 Computational Complexity